



INTERNSHIP REPORT

Computing Explicit Isomorphisms Between Quaternion Algebras: Algorithms and Elliptic Curve Applications

Tommy Chakroun

Supervised by

Travis MORRISON

June 2025

Abstract

In number theory, quaternion algebras are special algebras of dimension 4. They appear in pure number theory in the study of ternary quadratic forms. They also appear in elliptic curve theory as the endomorphism ring of an elliptic curve extended to the rationals. For reasons detailed in Section 3, the algorithmic problem of determining whether two quaternion algebras are isomorphic, and if so computing an explicit isomorphism, is interesting.

En théorie des nombres, les algèbres de quaternions sont des algèbres particulières de dimension 4. Elles apparaissent en théorie des nombres pure dans l'étude des formes quadratiques ternaires. Elles interviennent également dans la théorie des courbes elliptiques comme anneaux d'endomorphismes d'une courbe elliptique étendus aux rationnels. Pour des raisons détaillées dans la Section 3, le problème algorithmique consistant à déterminer si deux algèbres de quaternions sont isomorphes, et le cas échéant à calculer un isomorphisme explicite, présente un intérêt.

Contents

1	Introduction	4
2	Quaternion Algebras	7
3	Interlude: Motivation from Elliptic Curve Theory	10
4	Isomorphism via Ternary Quadratic Forms	11
4.1	Isotropic Quadratic Forms	11
4.2	Isomorphism between Split Quaternion Algebras	13
4.3	Isomorphism between Quaternion Algebras	14
5	Isomorphism via Splitting Matrix Algebras	17
6	Splitting a Matrix Algebra	20
6.1	Reduction to Finding Zero Divisors and Idempotents	20
6.2	The case $M_4(\mathbb{Q})$	22
7	Orders in Central Simple Algebras	23
7.1	Orders in central simple algebras	23
7.2	Discriminant	25
7.3	Computing Maximal Orders	27
8	Algorithms for Finding Zero Divisors	31
8.1	Use of randomization depending on the ground field	31
8.2	Heuristic Lifting of Idempotents over \mathbb{Q}	32
A	Central Simple Algebras	34
B	Implementation and Precise Algorithms	38
B.1	Reduction $\varphi : A \otimes_{\mathbb{Q}} B^{op} \rightarrow M_4(\mathbb{Q})$ to $f : A \rightarrow B$	38

B.2	Solution to a System $AX \equiv 0 \pmod{d}$ for Integral Matrices	40
B.3	Computation of Left Order	41
B.4	Randomized Computation of Central Idempotents over a Finite Field	42
C	Number of Irreducible Matrices in $M_n(\mathbb{F}_p)$	49

1. Introduction

Throughout this report, all fields are assumed to be commutative, and all algebras are assumed to be unital and associative.

Quaternion Algebras. Quaternion algebras are central simple algebras of dimension 4. When $\text{char}(F) \neq 2$, these algebras admit a very concrete description: they have a basis $1, i, j, ij$, with $i^2 = a \in F^\times$, $j^2 = b \in F^\times$, and $ij = -ji$. We denote such an algebra by $\left(\frac{a,b}{F}\right)$. See [14]. Quaternion algebras are closely related to ternary quadratic forms over F , a fact which underpins many computational methods.

Isomorphism of Quaternion Algebras: Statement of the Problem. The central computational task of this report is to solve the following problem.

Explicit Isomorphism of Quaternion Algebras

Let $A = \left(\frac{a,b}{\mathbb{Q}}\right)$ and $B = \left(\frac{c,d}{\mathbb{Q}}\right)$ be two quaternion algebras over \mathbb{Q} . Given that A and B are isomorphic as \mathbb{Q} -algebras, compute an explicit isomorphism $f : A \rightarrow B$. This means expressing the images of the basis elements of A under f in terms of the basis of B .

Motivation from Elliptic Curve Theory. Section 3 is dedicated to motivations coming from the theory of elliptic curves. Elliptic curves are special projective plane curves that carry a natural abelian group structure. The computational properties of these groups are of great interest in cryptography due to the hardness of the discrete logarithm problem. On a more abstract level, one can define *isogenies* between elliptic curves—morphisms of algebraic varieties that are also group homomorphisms. In particular, for an elliptic curve E , the set $\text{End}(E)$ of isogenies from E to itself forms a ring under composition. This endomorphism ring gives rise to computational problems with cryptographic applications. A surprising result is that for certain elliptic curves called *supersingular*, the extended endomorphism ring $\text{End}(E) \otimes_{\mathbb{Z}} \mathbb{Q}$ has the structure of a quaternion algebra over \mathbb{Q} .

Methods of this Report. Let $A = \left(\frac{a,b}{\mathbb{Q}}\right)$ and $B = \left(\frac{c,d}{\mathbb{Q}}\right)$ be two quaternion algebras over \mathbb{Q} that are assumed to be isomorphic. We investigate two main approaches to compute an explicit isomorphism $f : A \rightarrow B$.

The first method, a natural and direct approach, is to determine the algebraic conditions that the images of the basis elements of A must satisfy in B . This leads to a system of equations. We will see that this problem can always be reduced to

solving several equations of the type $aX^2 + bY^2 + cZ^2 = 0$ over \mathbb{Q} or a number field. This equation defines a (diagonal) ternary quadratic form. Fortunately, solving such equations is a classical problem, and efficient algorithms exist. We discuss the literature on this topic in more detail in subsection 4.1.

The second method is more conceptual and relies on the structure theory of central simple algebras. A key theorem states that $A \cong B$ if and only if the tensor product algebra $A \otimes_{\mathbb{Q}} B^{\text{op}}$ is a matrix algebra, specifically $A \otimes_{\mathbb{Q}} B^{\text{op}} \cong M_4\mathbb{Q}$. Letting $C = A \otimes_{\mathbb{Q}} B^{\text{op}}$, we can compute its structure constants for a basis e_1, \dots, e_{16} where $e_i e_j = \sum_k c_{ijk} e_k$. The problem is thus reduced to the following more general, and well-studied, problem.

Explicit Isomorphism Problem

Given a \mathbb{Q} -algebra C by its structure constants, which is known to be isomorphic to $M_n(\mathbb{Q})$, compute an explicit isomorphism from C to $M_n(\mathbb{Q})$.

In the literature, the main strategy to solve this is to find special elements in C , such as zero divisors or rank-one elements (whose rank is independent of the choice of isomorphism to $M_n(\mathbb{Q})$). A significant advantage of this method is its connection to the arithmetic of the algebras. Given maximal orders in A and B , one can deduce a maximal order in C with minimal computation. The knowledge of such an order can then be leveraged to find the required special elements in C more efficiently.

Contents. Section 2 defines quaternion algebras and states their basic properties, including different definitions, the standard involution, and the correspondence with ternary quadratic forms. Section 3 provides the motivation for computing isomorphisms of quaternion algebras from the theory of elliptic curves. Section 4 details the first method, which relies on solving ternary quadratic forms. Section 5 presents the reduction for the second main approach: reducing the problem to finding an isomorphism $A \otimes B^{\text{op}} \cong M_4\mathbb{Q}$. Section 6 tackles the problem of splitting a matrix ring by reducing it to finding special elements within the algebra, suggesting the use of orders, and summarizing the literature. Section 7 covers the necessary preliminaries about orders in central simple algebras, concluding with an algorithm to compute a maximal order, particularly for an algebra isomorphic to $M_n\mathbb{Q}$. Section 8 describes a method to find zero divisors in a matrix algebra by reduction to a finite field and then lifting the result, a potentially heuristic process. Appendix A provides foundational material on central simple algebras. Appendix B gives more precise algorithms and implementation details. Appendix C presents a formula for the number of matrices in $M_n\mathbb{F}_p$ with an irreducible minimal polynomial, used to justify certain randomized algorithms.

Implementation. An implementation of most of the algorithms presented in this report, and more, is available on GitHub:

`https://github.com/TommyChakroun/quat_alg_project`

In particular, we provide functions to: compute an isomorphism by solving ternary quadratic forms; compute a maximal order in a central simple algebra over \mathbb{Q} ; deduce an isomorphism from an explicit isomorphism $A \otimes B^{\text{op}} \rightarrow \text{M4}\mathbb{Q}$; and attempt to compute rank-one elements in an algebra isomorphic to $\text{M4}\mathbb{Q}$.

2. Quaternion Algebras

Definition. Let F be a field of characteristic not 2. For all $a, b \in F^\times$, there exists a unique unitary and associative F -algebra B , up to F -algebra isomorphism, with two elements $i, j \in B$ such that $1, i, j, ij$ form an F -basis and satisfy the relations:

$$i^2 = a, \quad j^2 = b, \quad \text{and} \quad ij = -ji. \quad (1)$$

We denote such an algebra by

$$\left(\frac{a, b}{F} \right) = F \oplus Fi \oplus Fj \oplus Fij.$$

These are called *quaternion algebras*.

The uniqueness up to F -algebra isomorphism is clear, and to check the existence, it suffices to write down the unique possible multiplication table which respects associativity and the relations (1) and then check that the whole table is well associative.

An important example: If F is a field of characteristic not 2, then $M_2(F) \simeq \left(\frac{1, 1}{F} \right)$ for the basis:

$$1 = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}, \quad i = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}, \quad j = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad ij = \begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix}.$$

Let $B = \left(\frac{a, b}{F} \right)$ be a quaternion algebra over a field F of characteristic not 2, and denote $1, i, j, ij$ an adapted basis. The sub- F -vector space H spanned by i, j, ij is exactly the set of the elements $x \in B$ such that $x \notin F$ and $x^2 \in F$. Thus, H is a supplement of F in B which does not depend on the choice of the basis. We define the conjugate of an element $s \in B$ to be the symmetric in the direction of H parallel to F :

$$\bar{s} := \lambda - h.$$

The conjugation on B enjoys the following properties:

- Transition to the next section...

Standard involution.

Definition 2.1. Let F be a field and B be an F -algebra. An *involution* on B is an F -linear map $\bar{\cdot} : B \rightarrow B$ such that $\bar{\bar{1}} = 1$ and for all $x, y \in B$, $\bar{\bar{x}} = x$ and $\overline{xy} = \bar{y}\bar{x}$. An involution is *standard* if for all $x \in B$, $x\bar{x} \in F$.

If B has a standard involution, then we also have $x + \bar{x} \in F$ for all $x \in B$. Indeed, $(1+x)\overline{(1+x)} \in F$ and we can write

$$(1+x)\overline{(1+x)} = (1+x)(1+\bar{x}) = 1 + x + \bar{x} + x\bar{x}.$$

We also have $x\bar{x} = \bar{x}x$.

We call $\text{trd}(x) = x + \bar{x}$ the *reduced trace* of x and $\text{nrd}(x) = x\bar{x}$ the *reduced norm* of x .

To check that an involution is standard, it is not enough to check that $x\bar{x} \in F$ holds for x in a basis of B . But we have the following lemma.

Lemma 2.2. *An involution over a finite-dimensional algebra B is standard if and only if there exists a basis e_1, \dots, e_n of B such that for all $1 \leq i, j \leq n$, we have $e_i\bar{e}_i \in F$ and $(e_i + e_j)\overline{(e_i + e_j)} \in F$.*

Proof. The direct sense is immediate. For the converse, denote $n_i = e_i\bar{e}_i$ and $n_{i,j} = (e_i + e_j)\overline{(e_i + e_j)}$. Then for all $x = \sum x_i e_i$,

$$x\bar{x} = \sum x_i e_i \sum x_i \bar{e}_i = \sum x_i^2 n_i + \sum x_i x_j (e_i \bar{e}_j + e_j \bar{e}_i) = \sum x_i^2 n_i + \sum x_i x_j (n_{i,j} - n_i - n_j)$$

lies in F . □

Another remark is that if B has a standard involution, then every $x \in B$ satisfies the polynomial of degree two with coefficients in F :

$$T^2 - \text{trd}(x)T + \text{nrd}(x).$$

In particular, if B has a standard involution and e_1, \dots, e_n is a basis of B with $e_1 = 1$, then for $i \geq 2$, the coefficient of e_i in e_i^2 is $\text{trd}(e_i)$.

Algorithm 1: Has Standard Involution

Input: A finite-dimensional algebra B over a field F with a basis

$$e_1 = 1, e_2, \dots, e_n$$

Output: True if B has a standard involution, False otherwise.

Let $t(e_1) = 2$ and for $2 \leq i \leq n$, let $t(e_i) \in F$ be the coefficient of e_i in e_i^2 . Extend

t by F -linearity on B and for all $\xi \in B$, let $\bar{\xi} := t(\xi) - \xi$.

If $\bar{e}_i \bar{e}_j \neq \bar{e}_j \bar{e}_i$ for some $2 \leq i, j \leq n$, return False.

If $n_i := e_i \bar{e}_i$ or $n_{i,j} := (e_i + e_j)(\bar{e}_i + \bar{e}_j)$ is not in F for some $2 \leq i, j \leq n$, return

False; else, return True.

Proof of correctness. If B has a standard involution, the algorithm returns True. Conversely, if the algorithm returns True, then B has a standard involution. □

Quadratic forms.

Theorem 2.3. *Let F be a field of characteristic not 2. Then the map*

$$\left[\left(\frac{a, b}{F} \right) \right] \mapsto [-aX^2 - bY^2 + abZ^2]$$

induces a bijection between the classes of quaternion algebras over F up to F -algebra isomorphism and the classes of non-degenerate ternary quadratic forms over F up to similarity.

Proof. Well-defined: Let $B = \left(\frac{a, b}{F} \right)$ and $C = \left(\frac{c, d}{F} \right)$ be isomorphic via $f : B \rightarrow C$. The standard involution s is preserved, so $f(B^0) = C^0$, and f induces an F -linear isomorphism $u : F^3 \rightarrow F^3$ such that for $x, y, z \in F$,

$$f(xi + yj + zij) = u_1(x, y, z)i + u_2(x, y, z)j + u_3(x, y, z)ij.$$

Hence:

$$\begin{aligned} \text{nrd}(xi + yj + zij) &= -ax^2 - by^2 + abz^2, \\ &= \text{nrd}(f(xi + yj + zij)) \\ &= -cu_1^2 - du_2^2 + cd u_3^2. \end{aligned}$$

So the forms $-aX^2 - bY^2 + abZ^2$ and $-cX^2 - dY^2 + cdZ^2$ are equivalent.

Surjectivity: Let $Q \sim aX^2 + bY^2 + cZ^2$ with $a, b, c \in F^\times$. Then:

$$Q \sim abc(aX^2 + bY^2 + cZ^2) \sim (bc)X^2 + (ac)Y^2 + (ab)Z^2.$$

This form has discriminant $1 \in F^\times / F^{\times 2}$. Write it as $-\alpha X^2 - \beta Y^2 + \gamma Z^2$. Since $\alpha\beta\gamma$ is a square, $\frac{\alpha\beta}{\gamma}$ is a square, so the form is equivalent to $-\alpha X^2 - \beta Y^2 + \alpha\beta Z^2$.

Injectivity: Suppose $-aX^2 - bY^2 + abZ^2$ and $-cX^2 - dY^2 + cdZ^2$ are similar. Since they have the same discriminant... \square

Local-Global Principle. See appendix on Hasse-Minkowski theorem.

Theorem 2.4 (Local-Global Principle for Quaternion Algebras over \mathbb{Q}). *Let B be a quaternion algebra over \mathbb{Q} . Then B is determined up to isomorphism by the set of places v of \mathbb{Q} where B is ramified.*

More precisely, for each place v of \mathbb{Q} , let $B_v = B \otimes_{\mathbb{Q}} \mathbb{Q}_v$. Then B is uniquely determined (up to isomorphism) by the set

$$\text{Ram}(B) := \{v \text{ place of } \mathbb{Q} : B_v \text{ is a division algebra}\},$$

which is a finite set of even cardinality.

Hilbert symbol over \mathbb{Q}_p .

Application to the problem `IdentifyMatrixwRing`.

3. Interlude: Motivation from Elliptic Curve Theory

Short definition of elliptic curve.

Short definition of the endomorphism ring $\text{End}(E)$.

Short definition and existence of the dual endomorphism map $\text{End}(E) \rightarrow \text{End}(E), \alpha \mapsto \hat{\alpha}$

Implication that $\text{End}(E) \otimes_{\mathbb{Z}} \mathbb{Q}$ could be a quaternion algebra sometimes.

Short definition of super singular elliptic curve.

Why are they interesting?

Statement that if E is super singular then $\text{End}(E) \otimes_{\mathbb{Z}} \mathbb{Q}$ is well an quaternion algebra over \mathbb{Q}

If E is super singular over \mathbb{F}_p and ..., $\text{End}(E) \otimes_{\mathbb{Z}} \mathbb{Q}$ is exactly $B_{p,\infty}$.

Computational problem believed to be hard.

4. Isomorphism via Ternary Quadratic Forms

In this section, we present what is perhaps the most natural approach to explicitly compute an isomorphism between two quaternion algebras $A = \left(\frac{\alpha, \beta}{\mathbb{Q}}\right)$ and $B = \left(\frac{a, b}{\mathbb{Q}}\right)$.

If $f : A \rightarrow B$ is such an isomorphism, then $\mu = f(i)$ must be a pure quaternion in B such that $\mu^2 = \alpha$. Our first task is to find such an element μ , which involves solving a particular type of quadratic equation. After finding μ , we can find a pure quaternion $\nu \in B$ that anticommutes with μ (i.e., $\nu\mu = -\mu\nu$) and whose square is a non-zero rational number, say $\nu^2 = \gamma$.

It follows that $\{1, \mu, \nu, \mu\nu\}$ is a \mathbb{Q} -basis for B satisfying $\mu^2 = \alpha$ and $\nu^2 = \gamma$. However, we may have $\gamma \neq \beta$. To conclude, we must find a new element $\nu' = (x + y\mu)\nu$ for some $x, y \in \mathbb{Q}$, such that $(\nu')^2 = \beta$. This is always possible and once again involves solving a quadratic equation.

This section is organized as follows. We begin by defining the type of quadratic equations we will encounter and review the state-of-the-art algorithms for solving them. For convenience, we then present the algorithm for the split case before addressing the general case for non-split algebras.

4.1. Isotropic Quadratic Forms. For motivation, we note that the subsequent sections will require solving equations of the following forms:

- $aX^2 + bY^2 - abZ^2 = \alpha$, with $a, b, \alpha \in \mathbb{Q}$ and X, Y, Z unknown in \mathbb{Q} .
- $U^2 - \alpha V^2 = b$, with U, V in a quadratic number field.
- $X^2 - \alpha Y^2 = \beta/\gamma$, with $X, Y \in \mathbb{Q}$.

We unify these problems under a single definition and then review the state-of-the-art for the relevant algorithms.

Let K be a field of characteristic not equal to 2. A quadratic form of dimension $n \geq 1$ over K is a homogeneous polynomial $Q \in K[X_1, \dots, X_n]$ of degree 2. Any such form corresponds to a unique symmetric matrix $S \in M_n(K)$ such that $Q(X_1, \dots, X_n) = X^T S X$, where $X = (X_1, \dots, X_n)^T$. We define the determinant of Q , denoted $\det(Q)$, as $\det(S)$. The form Q is non-degenerate if $\det(Q) \neq 0$. The map from symmetric matrices to quadratic forms is a bijection. Likewise, the evaluation map from quadratic forms to functions $K^n \rightarrow K$ is injective.

Definition 4.1 (Quadratic Equation). Let K be a field of characteristic not equal to 2. We call a *quadratic equation of dimension $n \geq 1$* over K an equation of the form

$$Q(X_1, \dots, X_n) = b,$$

where Q is a quadratic form of dimension n over K and $b \in K^\times$. When $b = 0$, we call it an *isotropic quadratic equation*, and we seek a non-zero solution $(x_1, \dots, x_n) \in K^n$.

The literature on determining whether a quadratic equation has a solution, and if so, computing one, is well-studied. Let us start with some basic reductions.

First, any quadratic form Q can be efficiently transformed into a diagonal form using Lagrange's reduction algorithm. That is, in terms of matrices, if S is a symmetric matrix in $M_n(K)$, there exists an invertible matrix P (which can be computed efficiently) such that PSP^T is diagonal. However, as suggested in [12], this is not always the most effective strategy.

If Q is degenerate, i.e., $\det(Q) = 0$, diagonalization reduces it to a form in fewer than n variables. We can repeat this process until the form is non-degenerate. Hence, we can assume Q is non-degenerate.

If Q is anisotropic (i.e., $Q(v) = 0 \iff v = 0$), we can reduce the problem of solving $Q(x) = b$ to an isotropic problem by increasing the dimension by one, a process known as homogenization. Specifically, $Q(x_1, \dots, x_n) = b$ has a solution if and only if $Q(X_1, \dots, X_n) - bX_{n+1}^2 = 0$ has a non-zero solution. This reduction is effective: if (x_1, \dots, x_{n+1}) is a non-zero solution, then x_{n+1} must be non-zero (since Q is assumed anisotropic). We then obtain a solution to the original equation by dividing the coordinates by x_{n+1} .

If Q is non-degenerate and isotropic (i.e., $Q(v) = 0$ for some $v \neq 0$), one can solve $Q(x) = b$ for any $b \in K^\times$ using a non-zero solution v . The procedure is as follows: we view Q as a function on K^n , and let S be the symmetric matrix associated with Q . Since Q is non-degenerate, we can choose a vector u such that $v^T Su = 1/2$. We then define two new vectors: $e_1 = v$ and $e_2 = u - Q(u)v$. A direct calculation shows that $Q(e_1) = 0$, $Q(e_2) = 0$, and the bilinear form evaluates to $e_1^T Se_2 = 1/2$. It follows that for any $\lambda, \mu \in K$, we have $Q(\lambda e_1 + \mu e_2) = \lambda\mu$. Therefore, to find a solution for a given $b \in K^\times$, we can set $\lambda = b$ and $\mu = 1$. The vector $x = be_1 + e_2$ is a solution, since $Q(be_1 + e_2) = b$.

Hence, the problem is reduced to studying the *isotropic problem*.

State of the Art We present the state of the art for the *isotropic problem*: let K be a field of characteristic not 2, and let Q be a non-degenerate quadratic form of dimension n ($\det(Q) \neq 0$). The problem is to determine if $Q(X_1, \dots, X_n) = 0$ has a non-zero solution, and if so, to find one.

1. For $K = \mathbb{Q}$, in any dimension n , there are efficient algorithms to determine if a solution exists. This relies on the Hasse-Minkowski theorem, which involves

checking for solutions in \mathbb{Q}_p for each prime p dividing $2 \det(Q)$, and over \mathbb{R} . The existence of local solutions can be checked efficiently using the Hilbert symbol.

2. For $K = \mathbb{Q}$ and $n = 3$, D. Simon provides an efficient algorithm in [12].
3. For $K = \mathbb{Q}$ and $n \in \{4, 5\}$, Simon improves on his previous work to give an efficient solution in [10].
4. For K a number field and $n = 3$, the problem can be reduced to solving a norm equation. After diagonalization, we have an equation $aX^2 + bY^2 + cZ^2 = 0$, which can be transformed into an equation of the form $X'^2 - dY'^2 = e$. If d is a square in K , this is easily solved. Otherwise, letting $L = K(\sqrt{d})$, the problem becomes solving the norm equation $\mathcal{N}_{L/K}(X' + Y'\sqrt{d}) = e$. An algorithm for this is provided in [11]. This method is applicable when $K = \mathbb{Q}$, but the specialized algorithms mentioned previously are generally more efficient in that case.

4.2. Isomorphism between Split Quaternion Algebras. Let $A = \left(\frac{a,b}{\mathbb{Q}}\right)$ be a quaternion algebra over \mathbb{Q} . The following are equivalent:

- A is split, i.e., isomorphic to $M_2(\mathbb{Q})$.
- There exists a non-zero pure quaternion $e \in A$ such that $e^2 = 0$.
- The ternary quadratic form $aX^2 + bY^2 - abZ^2 = 0$ has a non-trivial rational solution.

Moreover, these equivalences are effective: given a solution to the quadratic form, one can efficiently construct an explicit isomorphism. The equivalence between (2) and (3) is clear: if $e = Xi + Yj + Zk$ is a pure quaternion, then its square is $e^2 = aX^2 + bY^2 - abZ^2$. If $\phi : M_2(\mathbb{Q}) \rightarrow A$ is an isomorphism, the image of a non-zero nilpotent matrix is a pure quaternion $e \in A$ with $e^2 = 0$. Conversely, if such an e exists, $V = Ae$ is a two-dimensional right ideal of A . The map from A to $\text{End}_{\mathbb{Q}}(V)$ given by left multiplication is an isomorphism, and choosing a basis for V yields an explicit isomorphism $A \cong M_2(\mathbb{Q})$.

This yields the following algorithm.

Algorithm 2: Isomorphism from a split algebra to $M_2(\mathbb{Q})$

Input: A split quaternion algebra $A = \left(\frac{a,b}{\mathbb{Q}}\right)$.

Output: An isomorphism $f : A \rightarrow M_2(\mathbb{Q})$.

1. Find a non-trivial solution (X, Y, Z) to $aX^2 + bY^2 - abZ^2 = 0$;
 2. Define the non-zero element $e = Xi + Yj + Zk \in A$. Note that $e^2 = 0$;
 3. Find an element $v \in A$ such that $ev \neq 0$. For instance, choose v from $\{i, j\}$;
 4. Let the set $\mathcal{B} = \{ev, e(ev)\}$ be the basis for the right ideal Ae ;
 5. Construct the isomorphism $f : A \rightarrow M_2(\mathbb{Q})$ by mapping an element $u \in A$ to the matrix representing left multiplication by u with respect to the basis \mathcal{B} ;
-

The only computationally difficult step is the first one.

4.3. Isomorphism between Quaternion Algebras. Let $A = \left(\frac{\alpha, \beta}{\mathbb{Q}}\right)$ and $B = \left(\frac{a, b}{\mathbb{Q}}\right)$ be two isomorphic, non-split quaternion algebras over \mathbb{Q} , where a, b, α, β are integers.

Find μ such that $\mu^2 = \alpha$ First, we want to find an element $\mu \in B$ such that $\mu^2 = \alpha$. Since α is not a square in \mathbb{Q} , μ must be a pure quaternion.

First method: Writing $\mu = Xi + Yj + Zk$, we must solve the quadratic equation of dimension 3 over \mathbb{Q} : $aX^2 + bY^2 - abZ^2 = \alpha$. As the quadratic form on the left is anisotropic (since B is non-split), this is equivalent to finding a non-trivial solution to $aX^2 + bY^2 - abZ^2 - \alpha W^2 = 0$. For this, one can use the algorithm from [10]. This requires factoring the integer $(ab)^2\alpha$, but apart from this, the algorithm appears to be linear in the number of digits of a, b, α , according to the chart in [10].

Second method: The computer algebra system Magma uses an alternative method that is based on a trick but may be less efficient. Let $K = \mathbb{Q}(\sqrt{\alpha})$. Since $A \cong B$, their extensions of scalars A_K and B_K are also isomorphic. The algebra A_K is split, so $B_K = \left(\frac{a,b}{K}\right)$ must also be split. This implies we can find an isotropic vector for the quadratic form of dimension 3 over K : $U^2 - aV^2 - bW^2 = 0$, with $U, V, W \in K$. We can be sure that $W \neq 0$; otherwise, a would be a square in K .

Scaling the solution allows us to assume $W = 1$, so we find $U, V \in K$ such that $U^2 - aV^2 = b$. This means the element $\zeta = U + Vi + j \in B_K$ has reduced norm $\text{nrd}(\zeta) = U^2 - aV^2 - b = 0$. Now, write $U = u_1 + u_2\sqrt{\alpha}$ and $V = v_1 + v_2\sqrt{\alpha}$ for $u_i, v_i \in \mathbb{Q}$. We can express ζ as $\zeta_1 + \zeta_2\sqrt{\alpha}$, where $\zeta_1 = u_1 + v_1i + j$ and $\zeta_2 = u_2 + v_2i$ are elements of B . The element ζ_2 must be invertible; otherwise, we would have $U, V \in \mathbb{Q}$, which implies B is split, a contradiction. The condition $\text{nrd}(\zeta) = 0$ implies $\text{nrd}(\zeta_1\zeta_2^{-1} + \sqrt{\alpha}) = 0$. Let $\mu = \zeta_1\zeta_2^{-1}$. One can check from $U^2 - aV^2 = b$ that μ is a pure quaternion, and the norm condition becomes $\text{nrd}(\mu) + \alpha = 0$, which simplifies to $-\mu^2 + \alpha = 0$. Thus, $\mu^2 = \alpha$, as desired.

Overall, this method solves $U^2 - aV^2 - bW^2 = 0$ over $K = \mathbb{Q}(\sqrt{\alpha})$. The algorithm for this first reduces the form and then solves a norm equation over a quadratic extension of K using the algorithm from [11].

Adapting the final basis Next, we find a pure quaternion $\nu \in B$ that is orthogonal to μ (i.e., $\mu\nu = -\nu\mu$). This amounts to solving a system of linear equations for the coefficients of ν . Let its square be $\nu^2 = \gamma$.

So far, we have a basis $\{1, \mu, \nu, \mu\nu\}$ for B with $\mu^2 = \alpha$ and $\nu^2 = \gamma$. Since $A \cong B$, the element β/γ must be a norm from the extension $\mathbb{Q}(\sqrt{\alpha})/\mathbb{Q}$. We can therefore find $x, y \in \mathbb{Q}$ by solving the norm equation $X^2 - \alpha Y^2 = \beta/\gamma$. This is equivalent to finding an isotropic vector for the quadratic form $X^2 - \alpha Y^2 - (\beta/\gamma)Z^2 = 0$. For this, we use the algorithm from [12], specialized for dimension 3 over \mathbb{Q} . Any non-trivial solution will have $Z \neq 0$, so we can scale it to get a solution with $Z = 1$.

Finally, we define $\nu' = (x + y\mu)\nu$. This new element ν' still anticommutes with μ , and its square is $(\nu')^2 = (x + y\mu)\nu(x + y\mu)\nu = (x + y\mu)(x - y\mu)\nu^2 = (x^2 - \alpha y^2)\gamma = (\beta/\gamma)\gamma = \beta$. The isomorphism $f : A \rightarrow B$ is then given by $f(i) = \mu$ and $f(j) = \nu'$.

Conclusion Overall, to compute the isomorphism $f : A = \left(\frac{\alpha, \beta}{\mathbb{Q}}\right) \rightarrow B = \left(\frac{a, b}{\mathbb{Q}}\right)$, we have to solve:

1. $aX^2 + bY^2 - abZ^2 - \alpha W^2 = 0$ over \mathbb{Q} , or alternatively $U^2 - aV^2 - bW^2 = 0$ over the quadratic field $K = \mathbb{Q}(\sqrt{\alpha})$.
2. $X^2 - \alpha Y^2 - (\beta/\gamma)Z^2 = 0$ over \mathbb{Q} .

It seems preferable to solve equations entirely over \mathbb{Q} using the two algorithms of D. Simon, [12] and [10]. We can easily work with forms with integer coefficients. In these algorithms, the most computationally expensive step is factoring the determinant of the form. Thus, we only need to be able to factor the integers a, b, α, β , and the numerator and denominator of γ .

Performance

Here are the results from benchmarking the quaternion algebra isomorphism algorithm. The data is presented in two graphs: a combined performance and reliability plot, and a plot showing internal complexity growth.

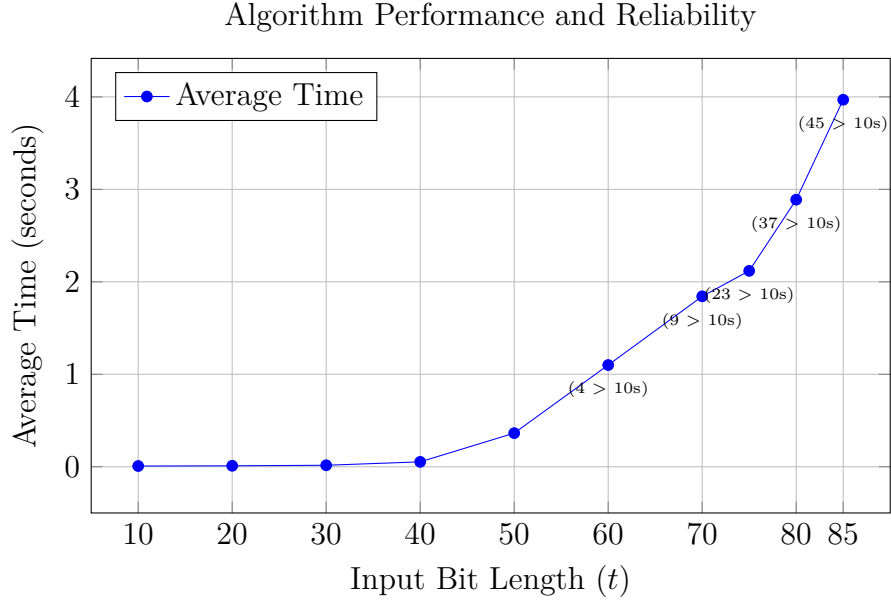


Figure 1: Average computation time as a function of the input bit length t . Labels below points indicate the number of tests (out of 100) that exceeded a 10-second timeout. The non-linear growth suggests an exponential complexity.

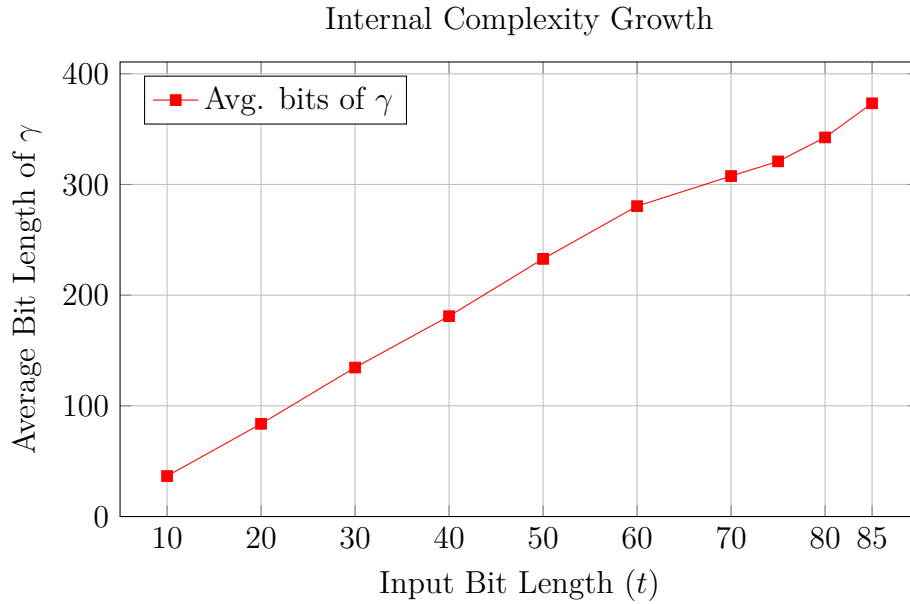


Figure 2: Average bit length of the intermediate value γ as a function of the input bit length t . The relationship appears to be strongly linear.

5. Isomorphism via Splitting Matrix Algebras

In this section, we present the second standard approach to explicitly computing an isomorphism between two quaternion algebras, $A = \left(\frac{\alpha, \beta}{\mathbb{Q}}\right)$ and $B = \left(\frac{a, b}{\mathbb{Q}}\right)$. This section presents the initial reduction steps, which use the theory of central simple algebras and the fact that quaternion algebras are precisely the central simple algebras of dimension 4. We refer to Appendix A for the necessary background on central simple algebras.

We will reduce our problem to a well-known problem in computational algebra: finding an explicit isomorphism for a matrix algebra, sometimes called "splitting a matrix algebra". This approach is interesting as it reveals certain arithmetic properties of our quaternion algebras, namely the notion of a maximal order, which will be detailed in a subsequent section. To my knowledge, this general problem has no efficient solution, even with randomization. The reader unfamiliar with maximal orders should read the relevant later sections of this report first. For the familiar reader: if we know a maximal order \mathcal{O}_1 in A and \mathcal{O}_2 in B , since the ramified places of A and B are easy to compute, we can deduce the factorization of the discriminants of \mathcal{O}_1 and \mathcal{O}_2 . It follows that $\mathcal{O}_1 \otimes \mathcal{O}_2$ is an order in $A \otimes B^{op}$ with a discriminant related to $(d_{\mathcal{O}_1} d_{\mathcal{O}_2})^{16}$, so we know its factorization. Computing a maximal order in $A \otimes B^{op}$ using the algorithm from [7] does not require integer factorization and runs in polynomial time (though it can be slow in practice). The knowledge of a maximal order in $A \otimes B^{op}$ could be crucial to solving the remaining problem (see, for example, [6]).

Lemma 5.1. *Let A be a central simple algebra over \mathbb{Q} of dimension n . There exists an isomorphism $\varphi : A \otimes A^{op} \rightarrow \text{End}_{\mathbb{Q}}(A)$ such that for all $a, b \in A$, $\varphi(a \otimes b)$ is the endomorphism $z \mapsto azb$. Consequently, since $\text{End}_{\mathbb{Q}}(A) \cong M_n(\mathbb{Q})$, we have $A \otimes A^{op} \cong M_n(\mathbb{Q})$.*

Proof. The map $A \times A^{op} \rightarrow \text{End}_{\mathbb{Q}}(A)$ which sends (a, b) to the endomorphism $z \mapsto azb$ is \mathbb{Q} -bilinear. By the universal property of the tensor product, this induces a \mathbb{Q} -linear map $\varphi : A \otimes A^{op} \rightarrow \text{End}_{\mathbb{Q}}(A)$ satisfying $\varphi(a \otimes b) = (z \mapsto azb)$ for all $a, b \in A$. It is easily checked that φ is an algebra homomorphism. Since $A \otimes A^{op}$ is a simple algebra, φ must be injective. By a dimension argument, it is also surjective, hence an isomorphism. \square

Proposition 5.2. *Let A and B be two central simple algebras over \mathbb{Q} of dimension n . Then A and B are isomorphic if and only if $A \otimes B^{op}$ is isomorphic to $M_n(\mathbb{Q})$. Moreover, this equivalence is explicit: the forward direction is computable in polynomial time, and the reverse direction is computable in randomized polynomial time.*

Proof. (\Rightarrow) Assume $A \cong B$. Let $f : A \rightarrow B$ be an isomorphism. Then $A \otimes B^{op} \cong B \otimes B^{op}$. By the previous lemma, $B \otimes B^{op} \cong M_n(\mathbb{Q})$. Explicitly, the map $\Phi : A \otimes B^{op} \rightarrow \text{End}_{\mathbb{Q}}(B) \cong M_n(\mathbb{Q})$, characterized on pure tensors by $\Phi(a \otimes b) = (z \mapsto f(a)zb)$, is an isomorphism. Given a basis \mathcal{B} for B , this map can be computed efficiently.

(\Leftarrow) For the reverse direction, we first prove the theoretical implication. The efficient randomized algorithm is presented in Appendix B, Section 6.1. The theoretical proof is not constructive, as it relies on the Wedderburn-Artin theorem. Assume $A \otimes B^{op} \cong M_n(\mathbb{Q})$. Tensoring both sides with B yields $(A \otimes B^{op}) \otimes B \cong M_n(\mathbb{Q}) \otimes B$. By associativity, this is equivalent to $A \otimes (B^{op} \otimes B) \cong B \otimes M_n(\mathbb{Q})$. From the previous lemma, $B \otimes B^{op} \cong M_n(\mathbb{Q})$, so we have $A \otimes M_n(\mathbb{Q}) \cong B \otimes M_n(\mathbb{Q})$. By the Wedderburn-Artin theorem, any central simple algebra is Brauer equivalent to its division algebra component. Since tensoring with a matrix algebra does not change the Brauer class, we have $[A] = [B]$ in the Brauer group. This means their division algebra components are isomorphic. As A and B have the same dimension, they must be isomorphic. \square

In the particular case of quaternion algebras, this proposition becomes:

Proposition 5.3. *Let A and B be two quaternion algebras over \mathbb{Q} . Then*

$$A \cong B \iff A \otimes B^{op} \cong M_4(\mathbb{Q}),$$

and this equivalence is algorithmically effective.

To consider $A \otimes B^{op}$ computationally and frame the problem in a standard form, we introduce the notion of an algebra given by *structure constants*.

Definition 5.4. An algebra C over \mathbb{Q} is given by **structure constants** if it is defined by a basis $\{e_1, \dots, e_N\}$ and rational numbers c_{ijk} for $i, j, k \in \{1, \dots, N\}$ such that, for all i, j , the product is given by

$$e_i e_j = \sum_{k=1}^N c_{ijk} e_k.$$

Given the quaternion algebras $A = \left(\frac{\alpha, \beta}{\mathbb{Q}}\right)$ and $B = \left(\frac{a, b}{\mathbb{Q}}\right)$, we can easily compute the structure constants for $C = A \otimes B^{op}$. We use the basis $\{e_1, \dots, e_{16}\}$ formed by the tensor product of the standard bases $\{1_A, i_A, j_A, k_A\}$ of A and $\{1_B, i_B, j_B, k_B\}$ of B^{op} . The structure constants c_{ijk} of C in this basis are expressed in terms of α, β, a, b , and many of them are zero.

This leads us to the following general problem:

Splitting $M_4(\mathbb{Q})$

Let C be an algebra over \mathbb{Q} given by structure constants, and assume $C \cong M_4(\mathbb{Q})$. Compute an explicit isomorphism $\varphi : C \rightarrow M_4(\mathbb{Q})$. That is, output 16 matrices in $M_4(\mathbb{Q})$ representing the images of the basis elements of C under φ .

This problem, which can be generalized to any dimension, is also referred to as the "splitting matrix algebra problem" or the "explicit isomorphism problem". It is well-studied in the literature, and the remainder of this report is dedicated to it.

6. Splitting a Matrix Algebra

In this section, we give a solution to the `ExplicitIsomorphismProblem`, with a special focus on the case where a \mathbb{Q} -algebra A is isomorphic to $M_4(\mathbb{Q})$. Our approach follows the methods in [6]. To construct a solution over \mathbb{Q} , we will first need to study the problem over other fields, such as a finite field \mathbb{F}_p or the field of real numbers \mathbb{R} , and for matrix algebras of smaller dimension ($n \leq 4$).

The central problem can be stated as follows:

ExplicitIsomorphismProblem

Given a central simple algebra A over a field F , presented by a basis, that is isomorphic to $M_n(F)$, compute an explicit isomorphism $\varphi : A \rightarrow M_n(F)$.

Our strategy is to first reduce the problem to the search for specific types of elements within the algebra. We then develop algorithms to find these elements over \mathbb{F}_p and \mathbb{R} . Finally, we address the main challenge of lifting these solutions from the simpler fields back to \mathbb{Q} .

6.1. Reduction to Finding Zero Divisors and Idempotents. Let F be a field, $n \geq 2$ be an integer, and A be an F -algebra of dimension $N = n^2$ given by an F -basis a_1, \dots, a_N and *structure constants* $(c_{i,j,k})_{1 \leq i,j,k \leq N}$. That is,

$$A = Fa_1 \oplus \dots \oplus Fa_N \quad \text{and} \quad a_i a_j = \sum_{k=1}^N c_{i,j,k} a_k \quad (1 \leq i, j \leq N)$$

with $c_{i,j,k} \in F$.

Assume that A is isomorphic to $M_n(F)$ as an F -algebra. We want to compute an explicit isomorphism $\varphi : A \rightarrow M_n(F)$, that is, to output a list of $N = n^2$ matrices with entries in F : M_1, \dots, M_N such that the unique F -linear map from A to $M_n(F)$ which maps a_i onto M_i is an F -algebra isomorphism.

First, we know that A inherits the following algebraic properties which hold in $M_n(F)$. The minimal polynomial of an element of A has a degree of at most n . An element in A is invertible if and only if it is invertible on the left or on the right. An element $x \in A$ is non-invertible if and only if it is a left zero divisor (there exists $y \in A$ such that $xy = 0$) if and only if it is a right zero divisor (there exists $y \in A$ such that $yx = 0$).

We can characterize the zero divisors in A more precisely. One can define the *rank* of an element $x \in A$ as follows. For every isomorphism $\varphi : A \rightarrow M_n(F)$, we have

$$n \cdot \text{rank}(\varphi(x)) = \dim_F(M_n(F)\varphi(x)) = \dim_F(\varphi(Ax)) = \dim_F(Ax).$$

Hence we define the *rank* of an element $x \in A$ by

$$\text{rank}_A(x) := \frac{1}{n} \dim_F(Ax) = \text{rank}(\varphi(x)) \quad (\text{for any isomorphism } \varphi : A \rightarrow M_n(F)).$$

In this section, we start by giving three efficient successive reductions of the problem based on finding some special kinds of elements in A . And then we will discuss which of these elements can be guessed by random picking, depending on the ground field F .

The first reduction is to seek a rank-one element in A . Finding an explicit isomorphism between A and $M_n(F)$ is equivalent to finding a rank-one element in A . Indeed, if we know $\varphi : A \rightarrow M_n(F)$, we just have to express a rank-one matrix, for example, the matrix $S \in M_n(F)$ with only 1 in the first row and 0 on other lines, in the basis $\varphi(a_1), \dots, \varphi(a_N)$: $S = \lambda_1 \varphi(a_1) + \dots + \lambda_N \varphi(a_N)$; then $x := \lambda_1 a_1 + \dots + \lambda_N a_N$ is a rank-one element. This involves solving an invertible linear system of size N , which is in $\mathcal{O}(n^6)$. Conversely, if $x \in A$ is a rank-one element, then Ax has dimension n over F , and $\varphi : A \rightarrow \text{End}_F(Ax), a \mapsto (z \mapsto az)$ is a non-zero F -algebra homomorphism, so it is injective because A is simple as a ring and hence bijective by a dimension argument. We start by building a basis \mathcal{B} of Ax ; this consists of extracting a linearly independent family from a generating set, which can be done in $\mathcal{O}(n^5)$. Then, for each $1 \leq i \leq N$, we write down the matrix of the map $Ax \rightarrow Ax, z \mapsto a_i z$ in the basis \mathcal{B} . This can be done in $\mathcal{O}(Nn^2) = \mathcal{O}(n^4)$. Hence we get an explicit isomorphism.

The second reduction is to seek a zero divisor in A , that is, an element with rank $0 < r < n$. We claim that if we can find a zero divisor in algebras isomorphic to $M_k(F)$ for all $2 \leq k < n$, then we can find a rank-one element in A . We proceed by induction on n . For $n = 2$, it is clear. Assume the reduction holds for all $k < n$. Let $x \in A$ be a zero divisor and let $r < n$ be the rank of x . By choosing an isomorphism $\varphi : A \rightarrow M_n(F)$, $X := \varphi(x)$ has rank r , so there exist P, Q invertible matrices such that $X = PJ_r Q$, where

$$J_r := \text{diag}(1, \dots, 1, 0, \dots, 0).$$

We see that $XZX = X$ has a solution $Z \in M_n(F)$, namely $Z = Q^{-1}J_r P^{-1}$. Hence the equation $xzx = x$ has a solution $z \in A$, and we can compute one solution by solving a system of N^2 linear equations in N^2 variables. Let $e := zx \in A$; e satisfies $e^2 = e$ and has rank r . Then $B := eAe$ is a subalgebra of A with unit e . One can easily compute a basis of B and the structure constants of B in this basis. We have $B \cong M_r(F)$: to show this, again we use our isomorphism φ . Let $E := \varphi(e) \in M_n(F)$. E satisfies $E^2 = E$ and has rank $r < n$. By linear algebra, there exists an invertible matrix P such that $E = PJ_r P^{-1}$. Then as F -algebras, we have the isomorphisms $B \cong EM_n(F)E \cong PJ_r P^{-1}M_n(F)PJ_r P^{-1} \cong J_r M_n(F)J_r \cong M_r(F)$. Hence we can apply our induction hypothesis to build $y \in B$, a rank-one element in B . By viewing y in the isomorphism to $M_r(F)$, we see that $\text{rank}_A(y) = \text{rank}_B(y) = 1$.

The third reduction is to find an element with a reducible minimal polynomial. Indeed, if $x \in A$ has a reducible minimal polynomial $\pi = fg$ in $F[T]$, then $f(x)$ is a zero divisor in A . Note that computing the minimal polynomial of x can be done by computing the matrix of the map $z \mapsto xz$ and then computing its minimal polynomial. The complexity of this task depends on the field F and the available algorithms to factorize polynomials in $F[T]$.

The fourth, very concrete reduction that we will try to use the most is to find a *non-trivial idempotent* in A , that is, $e \in A$ such that $e \neq 0, 1$ and $e^2 = e$.

6.2. The case $M_4(\mathbb{Q})$. To solve the explicit isomorphism problem in $M_4(\mathbb{Q})$, we only need to be able to find a zero divisor in $M_4(\mathbb{Q})$ and $M_2(\mathbb{Q})$. Indeed, if $A \cong M_4(\mathbb{Q})$ and $x \in A$ is a zero divisor, then the rank r of x is $r = 1, 2$, or 3 . If $r = 1$, we are done. If $r = 3$, we can compute y such that $xy = 0$. For that, we compute $\pi(T)$, the minimal polynomial of x over \mathbb{Q} . We can write $\pi(T) = Tg(T)$ and let $y = g(x)$. Hence y has rank one (thinking in the matrix space), and we are done. If $r = 2$, we compute $e \in A$ of rank 2 such that $e^2 = e$, we compute the algebra $B := eAe$ which is isomorphic to $M_2(\mathbb{Q})$, and then a zero divisor in B is a rank-one element in A and we are done.

The problem to find a zero divisor in $A \cong M_2(\mathbb{Q})$ is computationally equivalent to finding a non-trivial rational solution to a ternary quadratic form of the type

$$-aX^2 - bY^2 + abZ^2 = 0.$$

There exist some algorithms for this in the literature: statement: [13, Theorem 7.19], algorithm: [1, Algorithm I], [8, Theorem 3], [12, Algorithm 3.4].

7. Orders in Central Simple Algebras

7.1. Orders in central simple algebras. In this section, we fix R to be a principal ideal domain, F its field of fractions, and B a central simple algebra over F . (We will apply the results of this section with $R = \mathbb{Z}$ or $R = \mathbb{Z}_{(p)}$ for some prime p , so $F = \mathbb{Q}$ always.)

The assumption that R is a principal ideal domain allows us to give a simpler definition of an order. For more general cases over different rings, see [14], Chapters 9 and 10.

An R -order in B is a subring $\Lambda \subset B$ such that:

- Λ is an R -full lattice in B ; that is, there exists an F -basis e_1, \dots, e_N of B such that $\Lambda = Re_1 \oplus \dots \oplus Re_N$.
- Λ is a subring of B ; in particular, $1 \in \Lambda$.

In particular, Λ is a free R -module with R -basis e_1, \dots, e_N .

Remark 7.1. If B is a central simple algebra over \mathbb{Q} , then we can view \mathbb{Q} as the field of fractions of \mathbb{Z} or of $\mathbb{Z}_{(p)}$ for some prime p . In this context, there is a meaningful distinction between a \mathbb{Z} -order and a $\mathbb{Z}_{(p)}$ -order in B .

Theorem 7.2. *If Λ is an R -order in B , then for all $\alpha \in \Lambda$, α is integral over R ; hence $\text{trd}(\alpha) \in R$.*

Proof. Let $R[\alpha] := \{P(\alpha) \mid P \in R[X]\}$, which is an R -submodule of $\Lambda \cong R^n$. Since R is a principal ideal domain, it follows that $R[\alpha] \cong R^m$ for some integer $1 \leq m \leq n$. By choosing an R -basis of $R[\alpha]$, the multiplication-by- α map

$$R[\alpha] \rightarrow R[\alpha], \quad z \mapsto \alpha z$$

is represented by some matrix $M \in M_m(R)$ in this basis. Hence the characteristic polynomial of M is a monic polynomial in $R[X]$ that vanishes at α , so α is integral.

The second assertion follows because R is integrally closed and by the same reasoning. \square

Definition 7.3. An R -order is *maximal* if it is maximal with respect to inclusion among the R -orders in B .

Let us now specialize to the case of most interest to us, namely when $F = \mathbb{Q}$ and $R = \mathbb{Z}$ or $\mathbb{Z}_{(p)}$. Let B be a central simple algebra over \mathbb{Q} , and let

$$\Lambda = \mathbb{Z}e_1 \oplus \cdots \oplus \mathbb{Z}e_N$$

be a \mathbb{Z} -order in B . We define its localization at a prime p as

$$\Lambda_{(p)} := \mathbb{Z}_{(p)}e_1 \oplus \cdots \oplus \mathbb{Z}_{(p)}e_N = \mathbb{Z}_{(p)}\Lambda.$$

The second expression justifies that this localization does not depend on the choice of \mathbb{Z} -basis. It is clear that $\Lambda_{(p)}$ is a $\mathbb{Z}_{(p)}$ -order in B .

We also observe that

$$\Lambda = \bigcap_{p \text{ prime}} \Lambda_{(p)},$$

which follows componentwise from $\mathbb{Z} = \bigcap_p \mathbb{Z}_{(p)}$.

The main result is the local-global correspondence:

Theorem 7.4. *Let Λ be a \mathbb{Z} -order in a central simple \mathbb{Q} -algebra B . Then Λ is a maximal \mathbb{Z} -order if and only if $\Lambda_{(p)}$ is a maximal $\mathbb{Z}_{(p)}$ -order for every prime p .*

Proof. The forward direction is straightforward. Suppose that for all primes p , the localized order $\Lambda_{(p)}$ is maximal. Now, assume $\Lambda \subset \Gamma$ for some \mathbb{Z} -order Γ . Then for all primes p , we have $\Lambda_{(p)} \subset \Gamma_{(p)}$, and by maximality of $\Lambda_{(p)}$, it follows that $\Lambda_{(p)} = \Gamma_{(p)}$. Hence, by the identity

$$\Lambda = \bigcap_p \Lambda_{(p)} = \bigcap_p \Gamma_{(p)} = \Gamma,$$

we conclude that $\Lambda = \Gamma$, so Λ is maximal.

Conversely, suppose Λ is maximal, but there exists a prime p and a $\mathbb{Z}_{(p)}$ -order $\mathcal{O}(p)$ in B such that

$$\Lambda_{(p)} \subsetneq \mathcal{O}(p).$$

For each prime $q \neq p$, define $\mathcal{O}(q) := \Lambda_{(q)}$, and set

$$\Gamma := \bigcap_q \mathcal{O}(q).$$

We claim that Γ is a \mathbb{Z} -order in B , and that $\Gamma_{(q)} = \mathcal{O}(q)$ for all primes q . Moreover, we will show that $\Lambda \subsetneq \Gamma$, which will contradict the maximality of Λ .

Let $\Lambda = \mathbb{Z}e_1 \oplus \cdots \oplus \mathbb{Z}e_N$ and $\mathcal{O}(p) = \mathbb{Z}_{(p)}f_1 \oplus \cdots \oplus \mathbb{Z}_{(p)}f_N$. Since both are $\mathbb{Z}_{(p)}$ -lattices of full rank, there exists an integer $s > 0$ such that

$$s\Lambda_{(p)} \subset \mathcal{O}(p) \subset s^{-1}\Lambda_{(p)}.$$

This same inclusion holds for all primes q , since for $q \neq p$, we have $\mathcal{O}(q) = \Lambda_{(q)}$, so trivially:

$$s\Lambda_{(q)} \subset \mathcal{O}(q) \subset s^{-1}\Lambda_{(q)}.$$

Taking intersections over all primes q , and using the identity $\Lambda = \bigcap_q \Lambda_{(q)}$, we obtain:

$$s\Lambda \subset \Gamma \subset s^{-1}\Lambda.$$

This shows that Γ lies between two scalar multiples of Λ , so Γ is also a full \mathbb{Z} -lattice in B . In particular, there exist elements $a_1, \dots, a_N \in B$ such that

$$\Gamma = \mathbb{Z}a_1 \oplus \dots \oplus \mathbb{Z}a_N.$$

Now consider the localization

$$\Gamma_{(q)} = \mathbb{Z}_{(q)}a_1 \oplus \dots \oplus \mathbb{Z}_{(q)}a_N.$$

We claim that $\Gamma_{(q)} = \mathcal{O}(q)$ for all primes q . Indeed, for each q , we have:

$$\Gamma_{(q)} = \left(\bigcap_r \mathcal{O}(r) \right)_{(q)} = \bigcap_r \mathcal{O}(r)_{(q)}.$$

But $\mathcal{O}(r)_{(q)} = B$ when $r \neq q$, so this intersection reduces to $\Gamma_{(q)} = \mathcal{O}(q)$, as desired.

Thus, Γ is a \mathbb{Z} -order in B , and since $\Gamma_{(p)} = \mathcal{O}(p) \supsetneq \Lambda_{(p)}$, we must have $\Gamma \supsetneq \Lambda$, contradicting the maximality of Λ . Therefore, no such $\mathcal{O}(p)$ can exist, and all $\Lambda_{(p)}$ must be maximal. \square

7.2. Discriminant. We use the same setup as in the previous section: R is a principal ideal domain, F its field of fractions, and B a central simple algebra over F . Denote $N := \dim_F(B)$.

The notion of the *discriminant* of an order is intended to measure inclusion relations between orders. In summary, we define the discriminant to satisfy the following properties:

1. For any R -order Λ in B , the discriminant d_Λ is a nonzero element of R , well-defined up to a unit, i.e., it corresponds to a unique nonzero ideal of R .
2. For any R -orders $\Lambda, \Gamma \subset B$, if $\Lambda \subset \Gamma$, then $d_\Gamma \mid d_\Lambda$.
3. If $\Lambda \subset \Gamma$ and $d_\Gamma = d_\Lambda$ up to a unit, then $\Lambda = \Gamma$.

Definition 7.5. Let $\Lambda = Re_1 \oplus \cdots \oplus Re_N$ be an R -order in B . The *discriminant* of Λ is defined as

$$d_\Lambda := \det(\text{trd}(e_i e_j))_{1 \leq i, j \leq N} \bmod R^\times.$$

Since for all $1 \leq i, j \leq N$, the element $e_i e_j$ lies in Λ , its reduced trace lies in R by Proposition 7.2. Hence, d_Λ is indeed an element of R , and it is nonzero since the bilinear form $(x, y) \mapsto \text{trd}(xy)$ is non-degenerate by Proposition A.5.

We are going to prove that the definition of the discriminant does not depend on the choice of the R -basis of Λ , and thereby verify properties 2 and 3 regarding inclusion.

Suppose Λ, Γ are two R -orders in B , given by R -bases:

$$\Lambda = Re_1 \oplus \cdots \oplus Re_N, \quad \Gamma = Rf_1 \oplus \cdots \oplus Rf_N.$$

Let $M \in M_N(F)$ be the matrix expressing the e_k in terms of the f_i , that is,

$$e_k = M_{1,k}f_1 + \cdots + M_{N,k}f_N \quad \text{for } 1 \leq k \leq N.$$

Then, for all $1 \leq i, j \leq N$,

$$\text{trd}(e_i e_j) = \text{trd} \left(\left(\sum_{k=1}^N M_{k,i} f_k \right) \left(\sum_{\ell=1}^N M_{\ell,j} f_\ell \right) \right) = \sum_{k,\ell=1}^N M_{k,i} M_{\ell,j} \text{trd}(f_k f_\ell).$$

In matrix form, this reads

$$(\text{trd}(e_i e_j)) = M^\top (\text{trd}(f_i f_j)) M,$$

and hence,

$$\det(\text{trd}(e_i e_j)) = \det(M)^2 \det(\text{trd}(f_i f_j)).$$

We deduce all the desired properties from this relation:

- If $\Lambda = \Gamma$, then M is invertible over $M_N(R)$, so $\det(M)^2 \in R^\times$, and the discriminant is well-defined up to a unit.
- If $\Lambda \subset \Gamma$, then M has entries in R , so $\det(M)^2 \in R$, and hence $d_\Gamma \mid d_\Lambda$.
- If $\Lambda \subset \Gamma$ and $d_\Gamma = d_\Lambda$ up to a unit, then $\det(M)^2 \in R^\times$. Since R is a domain, this implies $\det(M) \in R^\times$, so M is invertible over $M_N(R)$. Thus, the f_i are R -linear combinations of the e_j , so $f_i \in \Lambda$, and therefore $\Lambda = \Gamma$.

An important consequence is the following: if

$$\Lambda_1 \subset \Lambda_2 \subset \cdots \subset \Lambda_r$$

is a chain of strictly increasing R -orders, then the corresponding sequence of discriminants

$$d_r \mid d_{r-1} \mid \cdots \mid d_1$$

is a strictly decreasing sequence (with respect to divisibility) of elements in R . In particular, if

$$d_1 = \pi_1^{\alpha_1} \cdots \pi_s^{\alpha_s}$$

is the factorization of d_1 into irreducibles, then $r \leq \alpha_1 + \cdots + \alpha_s$.

This shows that there exists a maximal order containing Λ_1 , since R cannot have an infinite strictly descending chain with respect to divisibility.

So far, we have only used the fact that trd respects integrality, is F -linear, and that the bilinear form $(x, y) \mapsto \text{trd}(xy)$ is non-degenerate. While one could define a discriminant using another F -linear form with these properties, the choice of trd —which coincides with the usual trace when $B \cong M_n(F)$ —is justified by the following normalization:

Theorem 7.6. *Let Λ be an R -order in B . If d_Λ is a unit in R (i.e., $\text{disc}(\Lambda) = R$), then Λ is a maximal R -order. If $B \cong M_n(F)$ for some integer n , then the converse is also true.*

Proof. The direct implication is immediate.

Conversely, suppose $B \cong M_n(F)$. Then the converse is more subtle and will require further work... \square

In general, a maximal order does not necessarily have discriminant 1. However, the easy characterization of maximal orders in matrix algebras provides a concrete way to verify maximality, which is especially useful for our SageMath implementation in the case where the algebra is isomorphic to $M_4(\mathbb{Q})$.

7.3. Computing Maximal Orders. Let B be a finite-dimensional central simple algebra over \mathbb{Q} , and let Λ be a \mathbb{Z} -order in B given by a \mathbb{Z} -basis:

$$\Lambda = \mathbb{Z}e_1 \oplus \cdots \oplus \mathbb{Z}e_N.$$

We describe a method following [7] to construct a \mathbb{Z} -basis of a maximal order $\Gamma \supseteq \Lambda$.

For each prime p , define $\mathcal{A}_p := \Lambda/p\Lambda$, which is the quotient of Λ by the two-sided ideal $p\Lambda$. Then:

$$\mathcal{A}_p = \mathbb{F}_p e_1 \oplus \cdots \oplus \mathbb{F}_p e_N,$$

with the structure of a finite-dimensional \mathbb{F}_p -algebra of dimension $N = \dim_{\mathbb{Q}}(B)$. Concretely, the structure constants of \mathcal{A}_p can be computed by multiplying the basis elements $e_i e_j \in \Lambda$, writing the result in terms of the e_k , and reducing the coefficients modulo p .

We now consider the *Jacobson radical* of \mathcal{A}_p :

$$\text{Rad}(\mathcal{A}_p) := \{x \in \mathcal{A}_p \mid xM = 0 \text{ for all simple } \mathcal{A}_p\text{-modules } M\}.$$

Let

$$\mathcal{C}_p := \mathcal{A}_p / \text{Rad}(\mathcal{A}_p)$$

be the semisimple quotient algebra. Define the composition of natural projections:

$$\Phi_p : \Lambda \xrightarrow{\text{mod } p} \mathcal{A}_p \twoheadrightarrow \mathcal{C}_p.$$

Since Λ is a ring and a free \mathbb{Z} -module, \mathcal{C}_p inherits at least a \mathbb{Z} -module structure (not necessarily free). The map Φ_p is a homomorphism of \mathbb{Z} -modules.

We now state the main result:

Theorem 7.7 (Detection of Non-Maximality). *With the notation above, suppose $\Lambda \subset B$ is not a maximal order. Then there exists a prime $p \mid d_{\Lambda}$, and an ideal $\mathcal{K} \subset \mathcal{C}_p$, either the zero ideal or a minimal nonzero ideal, such that the preimage*

$$\mathcal{I} := \Phi_p^{-1}(\mathcal{K})$$

satisfies $O_L(\mathcal{I}) \supsetneq \Lambda$; that is, the left order of \mathcal{I} strictly contains Λ .

Proof. Since Λ is not maximal, Theorem 7.6 implies that $d_{\Lambda} \in \mathbb{Z}$ is not a unit and hence divisible by some prime p □

This theorem leads directly to an algorithm: given a \mathbb{Z} -order Λ in B , the above procedure either produces a strictly larger order or certifies that Λ is maximal.

Algorithm 3: Find a strictly bigger \mathbb{Z} -order ([7], Section 5)

Input: B -algebra over \mathbb{Q} of dimension N , given by structure constants;

Λ -order given by \mathbb{Z} -basis $\Lambda = \mathbb{Z}e_1 \oplus \cdots \oplus \mathbb{Z}e_N$.

Output: A \mathbb{Z} -basis of a strictly bigger order Γ , if one exists; otherwise, report that Λ is maximal.

1. Compute the discriminant $d_\Lambda = d(e_1, \dots, e_N) \in \mathbb{Z}$. Factor d_Λ .

foreach *prime p dividing d_Λ* **do**

2. Compute the \mathbb{F}_p -algebra $\mathcal{A} = \mathbb{F}_p e_1 \oplus \cdots \oplus \mathbb{F}_p e_N$ by structure constants and the projection $\Lambda \rightarrow \mathcal{A}$.

3. Compute a \mathbb{F}_p -basis of $\text{Rad}(\mathcal{A})$.

4. Compute the \mathbb{F}_p -algebra $\mathcal{C} = \mathcal{A} / \text{Rad}(\mathcal{A})$ and the projection $\phi : \mathcal{A} \rightarrow \mathcal{C}$.

5. Compute the total projection $\Phi : \Lambda \rightarrow \mathcal{C}$.

6. Compute a \mathbb{Z} -basis of $\mathcal{I}_0 = \ker(\Phi)$.

7. Compute a \mathbb{Z} -basis of $O_L(\mathcal{I}_0)$.

if $O_L(\mathcal{I}_0)$ *strictly contains* Λ **then**

 | **return** $\Gamma = O_L(\mathcal{I}_0)$

end

8. Compute the list of \mathbb{F}_p -bases of the minimal nonzero two-sided ideals $\mathcal{K} \subseteq \mathcal{C}$.

foreach \mathcal{K} *in the list* **do**

9. Compute the \mathbb{F}_p -algebra \mathcal{C}/\mathcal{K} by structure constant and the projection $\mathcal{C} \rightarrow \mathcal{C}/\mathcal{K}$.

10. Compute the total projection $\Phi_{\mathcal{K}} : \Lambda \rightarrow \mathcal{C}/\mathcal{K}$.

11. Compute a \mathbb{Z} -basis of $\mathcal{I} = \ker(\Phi_{\mathcal{K}})$.

12. Compute a \mathbb{Z} -basis of $O_L(\mathcal{I})$.

if $O_L(\mathcal{I})$ *strictly contains* Λ **then**

 | **return** $\Gamma = O_L(\mathcal{I})$

end

end

end

13. Return that Λ is maximal.

Comments.

- Computing the discriminant is relatively easy, but factoring it can be hard.
- To compute a basis of the radical, we use SageMath's internal implementation (`A.basis_radical()`), which is based on the algorithm in [4, Section 2.3.2] or [9, Section 2]. **Complexity:** ...
- Computing the quotient, projection, and composition is straightforward.

- To compute the kernel of a \mathbb{Z} -module homomorphism Φ from a free \mathbb{Z} -module of rank N to an \mathbb{F}_p -algebra of dimension r over \mathbb{F}_p , we choose a matrix $M \in M_{r,N}(\mathbb{Z})$ such that

$$\Phi(e_i) = M_{1,i}f_1 + \cdots + M_{r,i}f_r.$$

Then computing the kernel is equivalent to finding a \mathbb{Z} -basis of the solutions $X \in \mathbb{Z}^N$ to the equation $MX \equiv 0 \pmod{p}$.

- Computing a \mathbb{Z} -basis of a left order can be reduced to finding the integer solutions $X \in \mathbb{Z}^N$ to the equation $MX \equiv 0 \pmod{d}$, where M is an integer matrix of size $N^2 \times N$. This can be computationally expensive.
- If \mathcal{C} is a semisimple finite-dimensional algebra over the finite field \mathbb{F}_p , then its minimal nonzero ideals satisfy

$$\mathcal{C} = \mathcal{I}_1 \oplus \cdots \oplus \mathcal{I}_r.$$

To compute them, it suffices to find a family (e_1, \dots, e_r) of central elements (i.e., elements that commute with all elements of \mathcal{C}), which are idempotent ($e_i^2 = e_i$), pairwise orthogonal ($e_i e_j = 0$ for $i \neq j$), and such that $e_1 + \cdots + e_r = 1$, with (e_1, \dots, e_r) maximal for these conditions. Then take $\mathcal{I}_s = \mathcal{C}e_s = e_s\mathcal{C}$. This is implemented internally in Magma (`CentralIdempotent(C)`, Magma Handbook, Section 81.3.4), which uses the algorithm from [3, Section 3].

8. Algorithms for Finding Zero Divisors

8.1. Use of randomization depending on the ground field. Let $n \geq 2$ be an integer.

In $M_n(\mathbb{F}_p)$ Let p be a prime. We consider the ring $M_n(\mathbb{F}_p)$, equipped with the uniform probability distribution. Let M be a random variable uniformly distributed in $M_n(\mathbb{F}_p)$. Then

$$\mathbb{P}(M \text{ is a zero divisor}) = \frac{p^{n^2} - 1 - (p^n - 1)(p^n - p) \cdots (p^n - p^{n-1})}{p^{n^2}},$$

which tends to zero as $p \rightarrow \infty$. Thus, finding a zero divisor by random sampling is not guaranteed.

However, for polynomials over \mathbb{F}_p , we have:

$$\mathbb{P}(f \in \mathbb{F}_p[T] \text{ of degree } d \text{ is reducible}) = 1 - \frac{1}{dp^d} \sum_{k|d} \mu\left(\frac{d}{k}\right) p^k \approx 1 - \frac{1}{d},$$

where μ denotes the Möbius function.

More precisely, in Appendix C, we give an exact formula for the probability that a uniformly chosen matrix in $M_n(\mathbb{F}_p)$ has a reducible minimal polynomial. In particular, we show that

$$\mathbb{P}(M \in M_n(\mathbb{F}_p) \text{ has a reducible minimal polynomial}) \geq \frac{1}{2}$$

for all $n \geq 2$ and all primes p . We also discuss the asymptotic behavior of this probability. Here are the values for small n and p :

Table 1: Probability of a Reducible Minimal Polynomial

$p \setminus n$	2	3	4
2	0.7500	0.9023	0.9367
3	0.7407	0.8243	0.8729
5	0.6720	0.7542	0.8171

Overall, this probability is always relatively high (at least $\frac{1}{2}$). Therefore, by repeatedly sampling random matrices, one can efficiently find a matrix in $M_n(\mathbb{F}_p)$ with a reducible minimal polynomial. Consequently, the explicit isomorphism problem is relatively easy in $M_n(\mathbb{F}_p)$.

In $M_n(\mathbb{R})$ In $M_n(\mathbb{R})$, for the Lebesgue measure on \mathbb{R}^{n^2} , the set of zero divisors has measure 0. However, the probability that a real monic polynomial of degree d with coefficients uniformly distributed in $[-M, M]$ is irreducible is low (it is even 0 if d is odd). Consequently, the explicit isomorphism problem is "easy" in $M_n(\mathbb{R})$.

In $M_n(\mathbb{Q})$ In $M_n(\mathbb{Q})$, for any finite subset $S \subset \mathbb{Q}$, the probability that a matrix with entries in S has a reducible minimal polynomial is very low. Hence the explicit isomorphism problem in $M_n(\mathbb{Q})$ can't be efficiently solved by randomly picking elements.

8.2. Heuristic Lifting of Idempotents over \mathbb{Q} . Let A be a finite-dimensional \mathbb{Q} -algebra of dimension $N = n^2$ assumed to be isomorphic to $M_4(\mathbb{Q})$. Since we know how to compute a \mathbb{Z} -basis of a maximal order in A , we can view

$$A = \mathbb{Q}a_1 \oplus \cdots \oplus \mathbb{Q}a_N \cong M_4(\mathbb{Q})$$

such that the \mathbb{Z} -algebra

$$O = \mathbb{Z}a_1 \oplus \cdots \oplus \mathbb{Z}a_N$$

is a maximal order in A (and thus isomorphic to $M_n(\mathbb{Z})$) and the structure constants are $c_{i,j,k} \in \mathbb{Z}$:

$$a_i a_j = \sum_{k=1}^N c_{i,j,k} a_k \quad (1 \leq i, j \leq N).$$

The fact that $c_{i,j,k} \in \mathbb{Z}$ allows us to define for every field K , and especially finite fields, the K -algebra:

$$O_K := K a_1 \oplus \cdots \oplus K a_N$$

in the sense that O_K has a basis a_1, \dots, a_N with structure constants being the images of the integers $c_{i,j,k}$ in K .

We want to solve the explicit isomorphism problem for A . As we saw in Section 6.1, it suffices to find $e \in A$ such that $e \neq 0, 1$ and $e^2 = e$. One approach is the following: We choose a field K such that it is easy with randomization to find a non-trivial $f \in O_K$ such that $f^2 = f$, and then we try to lift f to an element in A .

With $K = \mathbb{R}$. In $O_{\mathbb{R}} := \mathbb{R}a_1 \oplus \cdots \oplus \mathbb{R}a_N$, we can randomly pick a $b = r_1 a_1 + \cdots + r_N a_N$ with coefficients in \mathbb{Q} such that the minimal polynomial of b over \mathbb{Q} , $\pi_{b,\mathbb{Q}}$, has a real root $\gamma \in \mathbb{R}$. Then $c := g(b)$ where $\pi_{b,\mathbb{Q}}(T) = (T - \gamma)g(T)$ is a zero divisor in $O_{\mathbb{R}}$, and we can deduce a non-trivial idempotent $f \in O_{\mathbb{R}}$ with coefficients in $\mathbb{Q}(\gamma)$. Concretely, we can write $f = x_1 a_1 + \cdots + x_N a_N$ with $x_i \in \mathbb{Q}(\gamma)$, and the relation $f^2 = f$ is equivalent to the relations:

$$\left(\sum_{1 \leq i, j \leq N} x_i x_j c_{i,j,k} \right) - x_k = 0 \quad (1 \leq k \leq N).$$

Hence by taking \tilde{x}_i , rational approximations of x_i , the element $e := \tilde{x}_1 a_1 + \cdots + \tilde{x}_N a_N \in A$ satisfies that $e^2 - e$ has small rational coefficients in the basis a_1, \dots, a_N .

With $K = \mathbb{F}_p$. In $O_{\mathbb{F}_p} := \mathbb{F}_p a_1 \oplus \cdots \oplus \mathbb{F}_p a_N$, we can easily, by using randomization, build $f \in O_{\mathbb{F}_p}$, a non-trivial idempotent. By writing $f = x_1 a_1 + \cdots + x_N a_N$ with $x_i \in \mathbb{F}_p$ and taking any lift $y_i \in \mathbb{Z}$ of $x_i \in \mathbb{F}_p$, the element $e := y_1 a_1 + \cdots + y_N a_N \in O \subset A$ satisfies $e^2 - e \equiv 0 \pmod{p}$, i.e., $e^2 - e$ has integer coefficients divisible by p in the basis a_1, \dots, a_N .

With $K = \mathbb{Q}_p$. By the previous paragraph, we can, in particular, build an element $e \in O \subset O_{\mathbb{Q}_p}$ such that $e^2 - e \equiv 0 \pmod{p}$. Then, using a version of Hensel's Lemma, as in [2] Lemma 5.1, if we define by induction $e_0 := e$ and $e_{k+1} := 3e_k^2 - 2e_k^3$, the sequence (e_k) satisfies $e_k^2 - e_k \equiv 0 \pmod{p^{2^k}}$, $e_{k+1} \equiv e_k \pmod{p^{2^k}}$, and $e_k \equiv e \pmod{p}$. Hence, the sequence (e_k) converges in $O_{\mathbb{Q}_p}$ to an element $f \in O_{\mathbb{Q}_p}$ such that $f^2 = f$. This element f is non-trivial because it is equal to e modulo p . The approximation e_k of f gives a non-trivial element of O that is quasi-idempotent, in the sense that the coefficients of $e_k^2 - e_k$ are divisible by p^{2^k} .

Provisional conclusion For the moment, we are able to build in A an element $e \neq 0, 1$ such that $e^2 - e$ has small coefficients in the basis a_1, \dots, a_N , either for the absolute value or for the p -adic absolute value. If we manage to find an element $e \in A$ such that $e^2 - e$ is exactly 0, we are done.

A. Central Simple Algebras

All algebras B over a field F are assumed to be unitary and associative, so that we can view $F \subset B$. We recall that the center of B is the set $Z(B)$ of elements $b \in B$ that commute with every $x \in B$. It is a sub- F -algebra of B .

An algebra is a *division algebra* if every non-zero element has a two-sided inverse. Note that the center of a division algebra is a field, and every division algebra has the structure of an algebra over its center.

An algebra B , or more generally a ring, is *simple* if its only two-sided ideals are $\{0\}$ and B . This is equivalent to saying that for any non-zero ring R , every ring homomorphism $B \rightarrow R$ is injective. Now we define an important class of algebras.

Definition A.1. Let F be a field. A *central simple algebra* over F is a finite-dimensional algebra over F whose center is exactly F and which has no non-trivial two-sided ideals.

The standard example of a central simple algebra is $B = M_n(F)$, called the *split* central simple algebra of dimension n^2 . The goal of this section is to show that, by extending the base field, every central simple algebra becomes isomorphic to a split one.

First, as simple algebras, central simple algebras enjoy the following strong structure theorem.

Wedderburn-Artin Theorem and Consequences.

Theorem A.2 (Wedderburn-Artin Theorem, weak version). *Let B be a finite-dimensional simple algebra over F . Then there exists a division algebra D over F and an integer $n \geq 1$ such that B is isomorphic to $M_n(D)$. Moreover, D is unique up to F -algebra isomorphism.*

The proof follows [5, Section 2.1] by Philippe Gille and Tamás Szamuely.

Proof. Let I be a minimal non-zero left ideal of B . Such an ideal exists because a left ideal is an F -subspace, so we can choose one of minimal positive dimension. Define the ring D to be the ring of B -linear endomorphisms of I :

$$D := \text{End}_B(I) = \{f : I \rightarrow I \mid f \text{ is a group homomorphism and } f(bx) = bf(x) \text{ for all } x \in I, b \in B\}.$$

This is a sub- F -algebra of $\text{End}_F(I)$, and we claim that it is a division algebra. Indeed, if $f : I \rightarrow I$ is a non-zero B -linear map, then its kernel and image are left ideals of B

contained in I . Since $f \neq 0$, $\ker(f) \neq I$ and $\text{im}(f) \neq \{0\}$. By the minimality of I , we must have $\ker(f) = \{0\}$ and $\text{im}(f) = I$. Therefore, f is an isomorphism and thus has an inverse in D .

Now, we consider the ring of D -linear endomorphisms of I :

$$\text{End}_D(I) = \{f : I \rightarrow I \mid f \text{ is a group homomorphism and } f \circ g = g \circ f \text{ for all } g \in D\}.$$

It is also a sub- F -algebra of $\text{End}_F(I)$.

For any $b \in B$, the map of left multiplication by b is in $\text{End}_D(I)$. This allows us to define a map:

$$\Phi : B \rightarrow \text{End}_D(I); \quad b \mapsto (x \mapsto bx).$$

It is immediate that Φ is an F -algebra homomorphism. Furthermore, $\ker(\Phi)$ is a two-sided ideal of B . Since Φ is non-zero (e.g., $\Phi(1) = \text{id}_I$), its kernel cannot be all of B . By the simplicity of B , $\ker(\Phi) = \{0\}$, so Φ is injective. Let us show surjectivity. First, note that the set

$$IB := \left\{ \sum_{i=1}^r x_i b_i \mid r \in \mathbb{Z}_{\geq 0}, x_i \in I, b_i \in B \right\}$$

is a non-zero two-sided ideal of B , and thus $IB = B$ by simplicity. Hence, we can write $1 = \sum_{i=1}^r x_i b_i$ for some integer $r > 0$. Now let $f \in \text{End}_D(I)$. We have:

$$f = f\Phi(1) = \sum_{i=1}^r f \circ \Phi(x_i) \circ \Phi(b_i).$$

For $i = 1, \dots, r$, and $z \in I$, since $\text{id}_I \cdot z$ is in $\text{End}_D(I)$, by definition f commutes with $\text{id}_I \cdot z$. This yields

$$f \circ \Phi(x_i)(z) = f(x_i z) = f(x_i)z = \Phi(f(x_i))(z),$$

hence $f \circ \Phi(x_i) = \Phi(f(x_i))$.

And finally,

$$f = \Phi\left(\sum f(x_i)b_i\right).$$

This concludes the proof of surjectivity. Thus, Φ realizes an F -algebra isomorphism between B and $\text{End}_D(I)$.

Finally, by the theory of modules over a division ring, I is a free D -module of some rank n , i.e., $I \cong D^n$. It follows that

$$\text{End}_D(I) \cong \text{End}_D(D^n) \cong M_n(D^{op}) \cong M_n(D)$$

as rings, and therefore as F -algebras. This completes the proof. \square

It follows that if K is an algebraically closed field, then every central simple algebra over K is isomorphic to some matrix algebra $M_n(K)$. Indeed if D is a division algebra over K then $D = K$: for all $x \in D$ the minimal polynomial of x over K is irreducible (because D is a domain) and so is of degree 1, hence $x \in K$.

Lemma A.3. *Let B be a finite dimensional algebra over F and K/F be a field extension. Then B is a central simple algebra over F if and only if $B \otimes_F K$ is a central simple algebra over K .*

Proof. See Philippe Gille and Tamás Szamuely [5], Lemma 2.2.2. □

Proposition A.4. *Let B be a finite dimensional algebra over F . Then there exists a field extension K/F such that $B \otimes_F K \cong M_n(K)$ as a K -algebra, for some integer n .*

Proof. It follows from taking K to be the algebraic closure of F in the previous lemma and the previous remark. □

Reduced trace. Let F be a field and let B be a central simple algebra over F . Let K be an algebraic closure of F so that $B \otimes_F K \cong M_n(K)$. Hence, we can embed B into $M_n(K)$ via

$$\varphi : B \rightarrow M_n(K), \quad b \mapsto b \otimes 1.$$

We define the *reduced trace* of $b \in B$ by

$$\text{trd}(b) := \text{trace}(\varphi(b)) \in K.$$

Proposition A.5. *Assume that F is of characteristic 0. The reduced trace of b does not depend on the choice of the isomorphism. It satisfies the following properties:*

1. *For all $b \in B$, we have*

$$\text{trd}(b) = \frac{1}{n} \text{trace}_F(B \rightarrow B, x \mapsto bx),$$

hence $\text{trd}(b) \in F$, and $\text{trd} : B \rightarrow F$ is an F -linear form.

2. *The map $B \times B \rightarrow F$, $(x, y) \mapsto \text{trd}(xy)$ is a non-degenerate F -bilinear form.*
3. *If $F = \text{Frac}(R)$ for some principal ideal domain R , and if $x \in B$ is integral over R , then $\text{trd}(x) \in R$.*

Lemma A.6 (Skolem-Noether Theorem for matrix rings). *Let K be a field. Every K -algebra automorphism of $M_n(K)$ is inner, that is, of the form $M \mapsto PMP^{-1}$ for some $P \in \text{GL}_n(K)$.*

Proof of Lemma A.6. Let $\phi : M_n(K) \rightarrow M_n(K)$ be a K -algebra automorphism. Let $I \subset M_n(K)$ be the left ideal of matrices with only the first column nonzero. Fix a nonzero $u_0 \in K^n$. Since ϕ is an automorphism, $\phi((u_0 \mid 0 \mid \cdots \mid 0)) \neq 0$. Thus, we can choose $x_0 \in K^n$ such that

$$\phi((u_0 \mid 0 \mid \cdots \mid 0))x_0 \neq 0.$$

Define the K -linear endomorphism

$$P : K^n \rightarrow K^n, \quad u \mapsto \phi((u \mid 0 \mid \cdots \mid 0))x_0.$$

For any $M \in M_n(K)$ and any $u \in K^n$, we have $PMu = \phi(M)Pu$, so $PM = \phi(M)P$. The kernel of P is stable under left multiplication by any matrix in $M_n(K)$. Since $Pu_0 \neq 0$, the kernel is not all of K^n , and since the only ideals of $M_n(K)$ are $\{0\}$ and $M_n(K)$, the kernel must be zero. Hence, P is an isomorphism, and

$$\phi(M) = PMP^{-1}.$$

□

Proof of Proposition A.5. The well-definedness follows directly from Lemma A.6. For any field K of characteristic 0 and any matrix $M \in M_n(K)$, we have

$$\text{trace}(M) = \frac{1}{n} \text{trace}_F(M_n(K) \rightarrow M_n(K), Z \mapsto MZ).$$

Hence,

$$\text{trd}(b) = \text{trace}_F(M_n(K) \rightarrow M_n(K), Z \mapsto \varphi(b \otimes 1)Z) = \text{trace}_F(B \otimes_F K \rightarrow B \otimes_F K, z \mapsto (b \otimes 1)z).$$

Since we can choose a K -basis of $B \otimes_F K$ of the form $e_1 \otimes 1, \dots, e_N \otimes 1$, where e_1, \dots, e_N is an F -basis of B , this trace equals that of the map $B \rightarrow B$, $z \mapsto bz$. This proves the first point, as F -linearity is evident.

Now let $x \in B$, and assume that $\text{trd}(xy) = 0$ for all $y \in B$. Then

$$\text{trace}(\varphi(xy \otimes 1)) = 0 \quad \text{for all } y \in B.$$

Hence,

$$\text{trace}(\varphi(x \otimes 1)\varphi(y \otimes 1)) = 0 \quad \text{for all } y \in B.$$

Since the set $\{y \otimes 1 \mid y \in B\}$ spans $B_K = B \otimes_F K$ over K , it follows that $\varphi(x \otimes 1)$ is orthogonal to all matrices in $M_n(K)$ under the trace pairing, and thus must be zero. Hence, $x \otimes 1 = 0$, and therefore $x = 0$. This proves the second point.

Finally, suppose that $x \in B$ is integral over R . Then the matrix $A := \varphi(x \otimes 1)$ is integral over R , so its eigenvalues lie in $K = F^{\text{alg}}$ and are integral over R . Therefore, the trace of A , being the sum of its eigenvalues, is also integral over R . But since we have already seen that $\text{trd}(x) = \text{trace}(A) \in F$, and R is integrally closed in F , it follows that $\text{trd}(x) \in R$. This proves the third point. □

B. Implementation and Precise Algorithms

We present here most of the harder algorithms used in the implementation: https://github.com/TommyChakroun/quat_alg_project

We give as well the mathematical idea and the algorithm together with a comment on the complexity or randomization if we used randomized algorithms.

B.1. Reduction $\varphi : A \otimes_{\mathbb{Q}} B^{op} \rightarrow M_4(\mathbb{Q})$ to $f : A \rightarrow B$. Let A, B be two central simple algebras of same dimension n . In Proposition 5.2 we showed the equivalence

$$A \cong B \iff A \otimes_{\mathbb{Q}} B \cong M_{n^2}(\mathbb{Q}).$$

The proof of the direct implication was already explicit and efficient, but for the converse we had argued theoretically with the Wedderburn decomposition. Here we present an efficient but randomized algorithm to show this converse implication.

Suppose we have $\varphi : A \otimes_{\mathbb{Q}} B^{op} \rightarrow M_4(\mathbb{Q})$ an isomorphism. Fix us e_1, \dots, e_N a basis of A and f_1, \dots, f_n a basis of B . Denote $N = n^2$ and $V = \mathbb{Q}^N$. For each $v \in V$ we consider:

$$\begin{aligned} \lambda_v : A &\rightarrow V, & a &\mapsto \varphi(a \otimes 1)v \\ \mu_v : B &\rightarrow V, & b &\mapsto \varphi(1 \otimes b)v. \end{aligned}$$

Both λ_v and μ_v are \mathbb{Q} -linear maps.

First we claim that if we find $v \in V$ such that both λ_v and μ_v are \mathbb{Q} -linear isomorphisms, then $f : A \rightarrow B, a \mapsto \mu_v^{-1}(\lambda_v(a))$ is a \mathbb{Q} -algebra isomorphism. Indeed f is a \mathbb{Q} -linear isomorphism, $\mu_v(1) = v = \lambda_v(1)$ so f maps 1_A to 1_B . For the preservation of the product: let $a, a' \in A$ and $b = f(a), b' = f(a')$ in B . Then

$$\begin{aligned} \mu_v(f(aa')) &= \varphi(aa' \otimes 1)v \\ &= \varphi(a \otimes 1)\varphi(a' \otimes 1)v \\ &= \varphi(a \otimes 1)\varphi(1 \otimes b')v \\ &= \varphi(1 \otimes b')\varphi(a \otimes 1)v \\ &= \varphi(1 \otimes b')\varphi(1 \otimes b)v \\ &= \varphi(1 \otimes bb')v \\ &= \mu_v(bb') \end{aligned}$$

This shows $f(aa') = f(a)f(a')$.

Hence it suffices to show the existence of a suitable $v \in V$. We do this for λ . This comes from the theoretical existence of an isomorphism from A to B as follows. Choose $f : A \rightarrow B$ an isomorphism and denote $\varphi_f : A \otimes_{\mathbb{Q}} B^{op} \rightarrow M_N(\mathbb{Q})$ the other isomorphism obtained. By Lemma A.6 there exists some invertible matrix $P \in M_N(\mathbb{Q})$ such that:

$$\forall c \in A \otimes_{\mathbb{Q}} B^{op}, \quad \varphi(c) = P\varphi_f(c)P^{-1}.$$

Then for $v \in V$ and $a \in A$ we have

$$P\lambda_{P^{-1}v}(a) = P\varphi(a \otimes 1)P^{-1}v = \varphi_f(a \otimes 1)v = \text{Mat}(b \mapsto f(a)b)v.$$

Hence if we take v to be the coordinates of 1_B in the basis of B we obtain:

$$P\lambda_{P^{-1}v} : A \rightarrow V, \quad a \mapsto \text{Mat}(f(a), B)$$

so it is an isomorphism.

Now from the existence of a suitable v we are going to deduce that almost all $v \in V$ satisfy this condition in the following sense.

Fix a basis e_1, \dots, e_N of A , and let $M_i = \varphi(e_i \otimes 1) \in M_N(\mathbb{Q})$. For all $v = (v_1, \dots, v_N) \in V$:

The matrix of λ_v in the basis e_1, \dots, e_N of A and the canonical basis of $V = \mathbb{Q}^N$ is, by columns:

$$\begin{pmatrix} M_1 v & \cdots & M_N v \end{pmatrix}.$$

Expanding, we have:

$$\begin{pmatrix} \sum_{j=1}^N m_{1j}^{(1)} v_j & \cdots & \sum_{j=1}^N m_{1j}^{(N)} v_j \\ \sum_{j=1}^N m_{2j}^{(1)} v_j & \cdots & \sum_{j=1}^N m_{2j}^{(N)} v_j \\ \vdots & \ddots & \vdots \\ \sum_{j=1}^N m_{Nj}^{(1)} v_j & \cdots & \sum_{j=1}^N m_{Nj}^{(N)} v_j \end{pmatrix}$$

where $M_i = (m_{kj}^{(i)})_{1 \leq k, j \leq N} \in M_N(\mathbb{Q})$.

So the determinant is of the form $P_A(v_1, \dots, v_N)$ for some $P_A \in \mathbb{Q}[T_1, \dots, T_N]$.

The fact that we found one suitable v proves that P_A is not identically zero. Then by the Schwartz-Zippel lemma for all finite subset $S \subset \mathbb{Q}$ we have $\text{Card}(Z(P_A) \cap S^N) \leq N \text{Card}(S)^{N-1}$.

Similarly for P_B associated to B and μ , $\text{Card}(Z(P_B) \cap S^N) \leq N \text{Card}(S)^{N-1}$.

Then a $v \in V$ satisfies both λ_v and μ_v are isomorphisms if and only if v is not a zero of P_A neither P_B . By taking $S \subset \mathbb{Q}$ sufficiently large such that $2N \text{Card}(S)^{N-1} < \text{Card}(S)^N$ this is possible. This concludes the proof and we deduce the following algorithm.

Algorithm 4: Isomorphism $f : A \rightarrow B$ from $\varphi : A \otimes B^{op} \rightarrow M_N(\mathbb{Q})$

Input: A, B two central simple algebras of dimension n and

$\varphi : A \otimes B^{op} \rightarrow M_N(\mathbb{Q})$ an isomorphism given in a basis $e_i \otimes f_j$ which is the tensor product of a basis of A and a basis of B

Output: $f : A \rightarrow B$ an isomorphism

1. Choose a finite subset $S \subset \mathbb{Q}$ such that $2N\text{Card}(S)^{N-1} < \text{Card}(S)^N$

foreach $v \in S^N$ **do**

2. Compute U the matrix of $A \rightarrow V$, $a \mapsto \varphi(a \otimes 1)v$ in the basis of A and the canonical basis of V

3. Compute M the matrix of $B \rightarrow V$, $b \mapsto \varphi(1 \otimes b)v$ in the basis of B and the canonical basis of V

4. **if** M and U are invertible **then**

5. Compute M^{-1}

6. **return** $f : A \rightarrow B$ given in the basis of A by:

$$f(e_i) = \sum_{j=1}^N M_{j,i}^{-1} U_{j,i} f_j$$

end

end

Complexity. Not evaluated yet.

B.2. Solution to a System $AX \equiv 0 \pmod{d}$ for Integral Matrices.

Algorithm 5: ModularMatrixKernel

Input: $A \in M_{m \times n}(\mathbb{Z})$, $d \in \mathbb{Z}_{>0}$

Output: $X_1, \dots, X_k \in \mathbb{Z}^n$ a \mathbb{Z} -basis of the \mathbb{Z} -submodule $\{X \in \mathbb{Z}^n \mid AX \equiv 0 \pmod{d}\}$

1. Compute the Smith normal form of A : $U \in \text{GL}_m(\mathbb{Z})$, $V \in \text{GL}_n(\mathbb{Z})$ such that $UAV = D$ is in diagonal form

2. If $t = \min(n, m)$, let Y_1, \dots, Y_t be defined by $Y_i = (0, \dots, 0, \frac{d}{\gcd(d, D_{i,i})}, 0, \dots, 0)$

3. If $n > m$, complete with $Y_j = (0, \dots, 0, 1, 0, \dots, 0)$ for $t < j \leq n$

4. **return** VY_1, \dots, VY_k

Implementation. The function `kernel_mod(A,d)` is available at [maximal_orders/maximal_orders_u](#)

Complexity. Not evaluated yet.

B.3. Computation of Left Order. Let B be a finite dimensional algebra over \mathbb{Q} of dimension N . Let

$$I = \mathbb{Z}e_1 \oplus \cdots \oplus \mathbb{Z}e_N$$

be a \mathbb{Z} -(full) lattice in B . The *left order* of I is the set

$$O_L(I) := \{\alpha \in B \mid \alpha I \subset I\}.$$

We claim that $O_L(I)$ is indeed an order. Indeed $O_L(I)$ is a \mathbb{Z} -submodule of B , is stable under multiplication, contains 1. It remains to show that $O_L(I)$ contains a \mathbb{Q} -basis of B . Starting from b_1, \dots, b_N any basis of B , then for all i there exists $t_i \in \mathbb{Z}_{>0}$ such that $t_i b_i \in O_L(I)$.

Now we want to compute explicitly a \mathbb{Z} -basis f_1, \dots, f_N of $O_L(I)$. First note that there exists $s \in \mathbb{Z}_{>0}$ such that $s \cdot 1_B \in I$. Hence $O_L(I) \cdot s \subset I$ so $O_L(I) \subset s^{-1}I$. After, for $\alpha \in s^{-1}I$ that we write $\alpha = s^{-1}(x_1 e_1 + \cdots + x_N e_N)$ with $x_i \in \mathbb{Z}$:

$$\alpha \in O_L(I) \iff \forall j, \alpha e_j \in I \iff \forall j, s^{-1}(x_1 e_1 e_j + \cdots + x_N e_N e_j) \in \mathbb{Z}e_1 \oplus \cdots \oplus \mathbb{Z}e_N$$

Let's write the *structure constants* of the basis e_1, \dots, e_N , that is:

$$e_i e_j = \sum_{k=1}^N c_{ijk} e_k \quad ((1 \leq i, j \leq N), c_{ijk} \in \mathbb{Q}).$$

Then

$$\alpha \in O_L(I) \iff \forall j, k, \sum_{i=1}^N s^{-1} c_{ijk} x_i \in \mathbb{Z}$$

or equivalently by denoting $X = (x_1, \dots, x_N) \in \mathbb{Z}^N$ and $T = (s^{-1} c_{ijk})$ the rational matrix with N^2 rows and N columns:

$$\alpha \in O_L(I) \iff TX \in \mathbb{Z}^{N^2}.$$

Finally by multiplying both sides by the least common multiple of the denominators of T this yields solving $AX \equiv 0 \pmod{d}$ for some integer matrix A of size $N^2 \times N$. Hence together with the previous algorithm we deduce a method to compute the left order of I .

Algorithm 6: LeftOrder

Input: B a finite dimensional algebra over \mathbb{Q} ; e_1, \dots, e_N a \mathbb{Q} -basis of B
representing $I = \mathbb{Z}e_1 \oplus \dots \oplus \mathbb{Z}e_N$

Output: A \mathbb{Z} -basis of the left order $O_L(I)$

1. Write $1_B = r_1e_1 + \dots + r_Ne_N$ with $r_i \in \mathbb{Q}$
 2. Let $s \in \mathbb{Z}_{>0}$ be the least common multiple of the denominators of r_i
 3. Compute $c_{i,j,k}$ the structure constants of the basis e_1, \dots, e_N , that is
 $e_ie_j = \sum c_{i,j,k}e_k$
 4. Let d be the lcm of the denominators of the $s \cdot c_{i,j,k}$ and let
 $A = (d \cdot s \cdot c_{i,j,k})_{1 \leq i \leq N; 1 \leq j, k \leq N} \in M_{N^2, N}(\mathbb{Z})$
 5. Let $X_1, \dots, X_\ell \in \mathbb{Z}^N$ be a \mathbb{Z} -basis of solutions of $AX \equiv 0 \pmod{d}$
 6. **return** $f_1, \dots, f_\ell \in B$ where $f_i = \frac{1}{s}(X_{i,1}e_1 + \dots + X_{i,N}e_N)$
-

Implementation. The function `left_order(B,Zbasis_I)` is available at [maximal_orders/maximal_orders_utilities.sage](#).

Complexity: In number of additions/multiplications. Computing structure constants: one inversion of $N \times N$ matrix ($\mathcal{O}(N^3)$), N^2 products of matrix-vector of size N ($\mathcal{O}(N^4)$). Solving the system $AX \equiv 0 \pmod{d}$ with A of size $N^2 \times N$ and X of size N ($\mathcal{O}(??)$). Total: $\mathcal{O}(??)$.

Remark B.1. If we suppose in addition that B is a division algebra or even just that every element of the \mathbb{Z} -basis e_1, \dots, e_N of the lattice $I = \mathbb{Z}e_1 \oplus \dots \oplus \mathbb{Z}e_N$ are invertible in B , then we can use a more efficient algorithm as follows. Let $\alpha \in B$, we have

$$\alpha \in O_L(I) \iff \forall j, \alpha e_j \in I \iff \forall j, \alpha \in e_j^{-1}I = \mathbb{Z}e_j^{-1}e_1 \oplus \dots \oplus \mathbb{Z}e_j^{-1}e_N$$

The last direct sum holds because $(e_j^{-1}e_1, \dots, e_j^{-1}e_N)$ is still a \mathbb{Q} -basis of B since $z \mapsto e_j^{-1}z$ is a \mathbb{Q} -linear automorphism of B .

Hence to compute the left order it suffices to compute an intersection of lattices. There are efficient algorithms for this based on the Hermite normal form. This is the method used in SageMath for the case of division quaternion algebras.

B.4. Randomized Computation of Central Idempotents over a Finite Field.

We begin with some general notions about rings. Let R be an arbitrary (unital) ring. There is a correspondence between:

1. The decomposition of R into ideals: $R = I_1 \oplus \dots \oplus I_r$,

2. The decomposition of 1_R into *central orthogonal idempotents* $1_R = e_1 + \cdots + e_r$, where $e_i \in Z(R)$, $e_i^2 = e_i$, and $e_i e_j = 0$ if $i \neq j$.

Now we turn to algebras over finite fields. Fix a finite field $F = \mathbb{F}_p$. If A is a *semisimple algebra* (i.e., an algebra isomorphic to a direct product of simple algebras), then by the Wedderburn–Artin theorem, there exist integers n_1, \dots, n_r and division algebras D_1, \dots, D_r over F such that

$$A \cong M_{n_1}(D_1) \times \cdots \times M_{n_r}(D_r).$$

Lemma B.2. *Consequently, A has exactly r minimal nonzero two-sided ideals I_1, \dots, I_r , which satisfy $A = I_1 \oplus \cdots \oplus I_r$. Moreover, writing $1 = e_1 + \cdots + e_r$ for the corresponding decomposition of the identity, the elements e_1, \dots, e_r are central orthogonal idempotents, and e_i is the identity of I_i .*

Proof. Let us fix an isomorphism as above. The image of $M_{n_i}(D_i)$ in A is simple, so its image I_i is a minimal nonzero two-sided ideal of A , and $A = I_1 \oplus \cdots \oplus I_r$. By the correspondence above, we obtain a decomposition of 1 into central orthogonal idempotents e_i , each acting as the identity on I_i .

Suppose J is another minimal nonzero two-sided ideal of A . Choose $x \in J$, $x \neq 0$, and write $x = x_1 + \cdots + x_r$ with $x_i \in I_i$. Since $x \neq 0$, assume $x_1 \neq 0$. Then $e_1 x = x_1 \neq 0$, so $I_1 \cap J \neq 0$. But both I_1 and J are minimal two-sided ideals, so $I_1 = J$. Hence, any minimal two-sided ideal must be one of the I_i . \square

Lemma B.3. *Conversely, if e_1, \dots, e_r are central orthogonal idempotents summing to 1, then the ideals Ae_i are exactly the minimal nonzero two-sided ideals of A , up to permutation.*

Proof. ... \square

Finding central orthogonal idempotents e_1, \dots, e_r summing to 1 in A reduces to finding them in the center $Z(A)$. We have:

$$Z(A) \cong Z(D_1) \times \cdots \times Z(D_r) \cong K_1 \times \cdots \times K_r,$$

where each K_i/F is a finite field extension.

So we are reduced to the following problem: let Z be a finite F -algebra isomorphic to a product

$$Z \cong K_1 \times \cdots \times K_r,$$

of finite field extensions of F . Find a maximal family e_1, \dots, e_r of central orthogonal idempotents in Z .

To solve this, we study the structure of F -algebras isomorphic to such a product. Let $a = (a_1, \dots, a_r) \in K_1 \times \dots \times K_r$. The minimal polynomial of a over F is

$$\pi_a = \text{lcm}(\pi_{a_1}, \dots, \pi_{a_r}),$$

which is square-free and can be written as $\pi_a = \pi_1 \dots \pi_s$, with $\{\pi_1, \dots, \pi_s\} = \{\pi_{a_1}, \dots, \pi_{a_r}\}$.

If $s \geq 2$, by the Chinese Remainder Theorem, there exist polynomials h_1, \dots, h_s such that

$$h_i \equiv 1 \pmod{\pi_i}, \quad h_i \equiv 0 \pmod{\pi_j} \text{ for } j \neq i.$$

Then for each i , $h_i(a)$ has the form:

$$w_i = h_i(a) = (h_i(a_1), \dots, h_i(a_r)) = (0, 0, \dots, 0, 1, 0, \dots, 0),$$

which gives a central orthogonal idempotent.

If $s = r$, then the w_i are exactly the elementary idempotents $(0, \dots, 0, 1, 0, \dots, 0)$, which form a maximal list of central orthogonal idempotents in $K_1 \times \dots \times K_r$.

This procedure can be applied in A since minimal polynomials and their factorization are abstract notions that can be computed from structure constants. So we randomly choose $a \in A$, compute its minimal polynomial π_a , factor it as $\pi_a = \pi_1 \dots \pi_s$, and attempt the construction above.

If $s \geq 2$, then

$$A = A\omega_1 \oplus \dots \oplus A\omega_s,$$

where each $A\omega_i$ is a unital subalgebra with identity $\omega_i = h_i(a)$. We can then recurse within each $A\omega_i$.

Note that the case $s = r$ may not occur when the prime p is small and r is large; for example, when $A \cong \mathbb{F}_p^r$.

Still, if $r \geq 2$, then $s \geq 2$ occurs for most $a \in A$, so a randomized approach works: pick random elements of A until the minimal polynomial splits. If no such element is found, then A is a field and we return 1_A .

We deduce a deterministic algorithm in theory, but the reasonable complexity is randomized.

Algorithm 7: CentralIdempotentsCommutativeSplit

Input: Z Semisimple commutative algebra over a finite field \mathbb{F}_p given by structure constants

Output: Primitive central idempotents $e_1, e_2, \dots, e_r \in A$

foreach a *in* Z **do**

1. Compute the minimal polynomial $\pi_a \in \mathbb{F}_p[T]$ of a .
 2. Factor $\pi_a = \pi_1 \dots \pi_s$ in $\mathbb{F}_p[T]$.
 3. **if** $s \geq 2$ **then**
 4. By the Chinese remainder theorem, compute $h_1, \dots, h_s \in \mathbb{F}_p[T]$ such that $h_i \equiv 1 \pmod{\pi_i}$, $h_i \equiv 0 \pmod{\pi_j}$ for $j \neq i$.
 5. Compute $\omega_i := h_i(a) \neq 0$ for $i = 1, \dots, s$.
 6. Compute a basis BA_i of $A_i := A\omega_i$.
 7. Compute the structure constants of A_i in BA_i and build the abstract version of A_i .
 8. $Res = []$;
 9. **foreach** $i = 1, \dots, s$ **do**
 10. Compute recursively
$$e_1, \dots, e_k = \text{CentralIdempotentsCommutative}(A_i).$$
 11. Lift e_1, \dots, e_k in Z and add it to res .
 - end**
 12. **return** Res
- end**

end

13. **return** the list 1_A

Implementation. The function `central_idempotent_commutative_split` is available at [minimal_ideals/minimal_ideals_manually.sage](https://github.com/ricardomartinb/mimosa/blob/master/minimal_ideals/minimal_ideals_manually.sage).

Comments. The algorithm is deterministic in theory with complexity at worst $\text{Card}(Z) = p^n$. It really depends on the size of the prime p and the dimension n . If both are relatively small we can easily iterate on Z . Otherwise if p is a very large prime around 100 bits then we can't reasonably iterate on Z so we prefer to pick a finite random number of elements of Z . Luckily when p is large and n small it happens that the probability that a is *decomposable* that is $s \geq 2$ in the algorithm is very high. Actually when p is very large in comparison to n we can use an easier algorithm. For example in our purpose of computing a maximal order containing an order O in the matrix ring $M_4(\mathbb{Q})$, for each prime dividing disc O we compute a finite dimensional algebra \mathcal{A} over \mathbb{F}_p of dimension 16 then we quotient \mathcal{A} by its radical so $\mathcal{A}/\text{Rad}(\mathcal{A})$ has dimension less than 16 over \mathbb{F}_p and consequently its center Z has dimension at most 16 over \mathbb{F}_p . It is in this final algebra Z that we run our algorithm.

Hence provided that all prime factors of disc O are much bigger than 16, the following algorithm may give the correct answer with a high probability and less computation time.

Algorithm 8: CentralIdempotentsCommutativeOneTime

Input: Z Semisimple commutative algebra over a finite field \mathbb{F}_p given by structure constants

Output: Primitive central idempotents $e_1, e_2, \dots, e_r \in A$

1. Pick a random $a \in Z$.
 2. Compute the minimal polynomial $\pi_a \in \mathbb{F}_p[T]$ of a .
 3. Factor $\pi_a = \pi_1 \dots \pi_s$ in $\mathbb{F}_p[T]$.
 4. By the Chinese remainder theorem, compute $h_1, \dots, h_s \in \mathbb{F}_p[T]$ such that $h_i \equiv 1 \pmod{\pi_i}$, $h_i \equiv 0 \pmod{\pi_j}$ for $j \neq i$.
 5. Compute $\omega_i := h_i(a) \neq 0$ for $i = 1, \dots, s$.
 6. **return** the list e_1, \dots, e_s
-

Implementation. The function `central_idempotent_commutative` is available at [minimal_ideals/minimal_ideals_manually.sage](#).

Precise evaluation of the probabilities. The success and complexity of the previous algorithm depend directly on the proportion of elements of Z satisfying good algebraic properties. To compute this proportion, we can as well work directly on

$$Z = K_1 \times \dots \times K_r$$

where K_i/\mathbb{F}_p are finite field extensions. Let us denote $s_i = [K_i : \mathbb{F}_p]$, so that $K_i \cong \mathbb{F}_{p^{s_i}}$. We denote also n the \mathbb{F}_p -dimension of Z , that is:

$$n = \sum_{i=1}^r s_i = [K_1 : \mathbb{F}_p] + \dots + [K_r : \mathbb{F}_p]$$

The total number of elements in the algebra is thus $\text{Card } Z = p^n$.

Since we will not know a priori the value of r when running the algorithm, n is the only reasonable bound that we can put on r (as $s_i \geq 1$, we have $r \leq n$).

Now let us define two subsets of Z :

$$S_0 := \{a = (a_1, \dots, a_r) \in Z \mid \pi_{a_1} = \dots = \pi_{a_r}\}$$

where π_x denotes the minimal polynomial of x over \mathbb{F}_p .

$$S_1 := \{a = (a_1, \dots, a_r) \in Z \mid \exists i \neq j, \pi_{a_i} = \pi_{a_j}\}$$

and let us denote $d_i = \frac{\text{Card } S_i}{\text{Card } Z}$ their densities: $0 \leq d_i \leq 1$.

Then:

1. The success of the algorithm `CentralIdempotentsCommutativeSplit` depends on whether d_0 is near zero for the first split, and afterward we have to deal with the recursion.
2. The success of the algorithm `CentralIdempotentsCommutativeOneTime` depends only on whether d_1 is near zero.

Computation of the density d_0 . Let $a = (a_1, \dots, a_r) \in S_0$. By definition, all components a_i share the same minimal polynomial over \mathbb{F}_p , let's call it $P(X)$. Let $\deg(P) = d$. For $P(X)$ to be the minimal polynomial of $a_i \in K_i = \mathbb{F}_{p^{s_i}}$, its degree d must be a divisor of s_i . This must hold for all $i = 1, \dots, r$. Therefore, the degree d of any such common minimal polynomial must divide $g = \gcd(s_1, \dots, s_r)$.

Let $N_d(p)$ be the number of monic irreducible polynomials of degree d over \mathbb{F}_p . For a fixed such polynomial $P(X)$ of degree d , the number of its roots in any field extension K_i where $d \mid s_i$ is exactly d . To form an element of S_0 corresponding to this polynomial $P(X)$, we must choose one of its d roots for each component a_i . This gives d^r choices.

Summing over all possible polynomials, we get the cardinality of S_0 :

$$\text{Card } S_0 = \sum_{d \mid g} N_d(p) \cdot d^r$$

To show that this quantity is small compared to $\text{Card } Z$, we can bound it. Since $d \leq g$ for any $d \mid g$, we have $d^r = d \cdot d^{r-1} \leq d \cdot g^{r-1}$.

$$\begin{aligned} \text{Card } S_0 &\leq \sum_{d \mid g} N_d(p) \cdot (d \cdot g^{r-1}) \\ &= g^{r-1} \sum_{d \mid g} N_d(p) \cdot d \end{aligned}$$

Using the classical identity $\sum_{d \mid m} N_d(p) \cdot d = p^m$, we obtain:

$$\text{Card } S_0 \leq g^{r-1} p^g$$

The density d_0 is therefore bounded by:

$$d_0 = \frac{\text{Card } S_0}{\text{Card } Z} \leq \frac{g^{r-1} p^g}{p^{\sum s_i}}$$

Since $g = \gcd(s_1, \dots, s_r) \leq s_i$ for all i , we have $\sum s_i \geq r \cdot g$. This gives:

$$d_0 \leq \frac{g^{r-1}p^g}{p^{rg}} = \frac{g^{r-1}}{p^{g(r-1)}} = \left(\frac{g}{p^g}\right)^{r-1}$$

For the algorithm to be in a non-trivial case, we must have $r > 1$. Since $p \geq 2$ and $g \geq 1$, the term p^g grows much faster than g , ensuring that $\frac{g}{p^g} < 1$. For instance, if $p = 2$ and $g = 1$, the ratio is $1/2$. If $g \geq 2$, the ratio is even smaller. Thus, for $r > 1$, the density d_0 is very small, confirming that picking an element from S_0 at random is a low-probability event. For example, if we take the smallest possible values $p = 2, r = 2, g = 1$, we get $d_0 \leq 1/2$. If $r = 3$, $d_0 \leq 1/4$, and so on. The probability decreases exponentially with r .

Computation of the density d_1 . (The mathematical analysis for this section will be completed later.)

C. Number of Irreducible Matrices in $M_n(\mathbb{F}_p)$

Let $n > 0$ be an integer and let p be a prime.

In this appendix, we give an explicit formula for the number $I(n, p)$ of *irreducible* matrices in $M_n(\mathbb{F}_p)$, that is, matrices with an irreducible minimal polynomial. This will allow us to calculate the probability that a uniformly random matrix in $M_n(\mathbb{F}_p)$ has a reducible minimal polynomial. If the minimal polynomial π_M of a matrix M is reducible, say $\pi_M = fg$ where f, g are non-constant polynomials, then $f(M)$ is a non-zero matrix (since $\deg(f) < \deg(\pi_M)$) which is a zero divisor because $f(M)g(M) = \pi_M(M) = 0$.

First, we recall and we will use the classical result: for an integer $d > 0$, the number of monic irreducible polynomials in $\mathbb{F}_p[X]$ of degree d is

$$N_d(p) = \frac{1}{d} \sum_{k|d} \mu\left(\frac{d}{k}\right) p^k$$

where μ is the Möbius function.

For $M \in M_n(\mathbb{F}_p)$, we denote by π_M its minimal polynomial. First, if $M \in M_n(\mathbb{F}_p)$ has a minimal polynomial π_M that is irreducible of degree d , then d must divide n . Indeed, let χ_M be the characteristic polynomial of M . We have that π_M divides χ_M . Since π_M is irreducible, the rational canonical form of M is composed of blocks of the companion matrix of π_M . This implies that $\chi_M = \pi_M^k$ for some integer k . By comparing degrees, we have $n = \deg(\chi_M) = k \cdot \deg(\pi_M) = kd$, and hence d divides n .

Conversely, let us fix $\pi \in \mathbb{F}_p[X]$ as a monic irreducible polynomial of degree $d \geq 1$ such that $d \mid n$. Let

$$\mathcal{A}_\pi := \{M \in M_n(\mathbb{F}_p) \mid \pi_M = \pi\}$$

and let us count the cardinal of \mathcal{A}_π . By the theory of rational canonical forms, a matrix M has the minimal polynomial π (with π irreducible of degree d) if and only if it is similar to a block diagonal matrix

$$D_\pi = \text{diag}(C_\pi, \dots, C_\pi)$$

where there are $k = n/d$ blocks, and C_π is the $d \times d$ companion matrix of π . If $\pi(X) = X^d + a_{d-1}X^{d-1} + \dots + a_0$, then

$$C_\pi = \begin{pmatrix} 0 & 0 & \dots & 0 & -a_0 \\ 1 & 0 & \dots & 0 & -a_1 \\ 0 & 1 & \dots & 0 & -a_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \dots & 1 & -a_{d-1} \end{pmatrix}$$

The set \mathcal{A}_π is the orbit of D_π under the conjugation action of $\text{GL}_n(\mathbb{F}_p)$ on $\text{M}_n(\mathbb{F}_p)$.

$$\mathcal{A}_\pi = \{PD_\pi P^{-1} \mid P \in \text{GL}_n(\mathbb{F}_p)\}$$

By the orbit-stabilizer theorem, the size of this orbit is

$$|\mathcal{A}_\pi| = \frac{|\text{GL}_n(\mathbb{F}_p)|}{|C_{\text{GL}_n(\mathbb{F}_p)}(D_\pi)|}$$

where $C_{\text{GL}_n(\mathbb{F}_p)}(D_\pi) = \{P \in \text{GL}_n(\mathbb{F}_p) \mid PD_\pi = D_\pi P\}$ is the centralizer of D_π in $\text{GL}_n(\mathbb{F}_p)$.

Let's compute the centralizer of D_π . A matrix $P \in \text{M}_n(\mathbb{F}_p)$, written as a $k \times k$ block matrix (P_{ij}) with blocks of size $d \times d$, commutes with D_π if and only if $P_{ij}C_\pi = C_\pi P_{ij}$ for all $i, j \in \{1, \dots, k\}$. The centralizer of a companion matrix C_π of an irreducible polynomial π is the ring of polynomials in C_π , denoted $\mathbb{F}_p[C_\pi]$. Since π is irreducible of degree d , this ring is a field isomorphic to \mathbb{F}_{p^d} . Thus, the centralizer of D_π in $\text{M}_n(\mathbb{F}_p)$, denoted $C_{\text{M}_n(\mathbb{F}_p)}(D_\pi)$, consists of block matrices where each block P_{ij} is in $\mathbb{F}_p[C_\pi]$. This ring is isomorphic to the matrix ring $\text{M}_k(\mathbb{F}_p[C_\pi])$, and thus to $\text{M}_k(\mathbb{F}_{p^d})$. The centralizer in $\text{GL}_n(\mathbb{F}_p)$ is the group of invertible matrices in this ring, which is isomorphic to $\text{GL}_k(\mathbb{F}_{p^d})$. Therefore, $|C_{\text{GL}_n(\mathbb{F}_p)}(D_\pi)| = |\text{GL}_k(\mathbb{F}_{p^d})|$, where $k = n/d$.

Finally, we have found that for a given irreducible polynomial π of degree $d|n$:

$$|\mathcal{A}_\pi| = \frac{|\text{GL}_n(\mathbb{F}_p)|}{|\text{GL}_{n/d}(\mathbb{F}_{p^d})|}$$

To conclude, the total number of irreducible matrices in $\text{M}_n(\mathbb{F}_p)$ is the sum of $|\mathcal{A}_\pi|$ over all possible irreducible polynomials π :

$$I(n, p) = \sum_{d|n} \sum_{\substack{\pi \text{ irreducible} \\ \deg(\pi)=d}} |\mathcal{A}_\pi| = \sum_{d|n} N_d(p) \cdot \frac{|\text{GL}_n(\mathbb{F}_p)|}{|\text{GL}_{n/d}(\mathbb{F}_{p^d})|}.$$

with

$$N_d(p) = \frac{1}{d} \sum_{k|d} \mu\left(\frac{d}{k}\right) p^k \quad \text{and} \quad |\text{GL}_m(\mathbb{F}_q)| = \prod_{i=0}^{m-1} (q^m - q^i)$$

This gives a complete, albeit complex, formula for $I(n, p)$. We compute this for small values of n and p .

Hence, the probability that a matrix has a reducible minimal polynomial is $1 - I(n, p)/p^{n^2}$, which is given in the following table for small n, p :

Table 2: Counts and Probabilities of Irreducible Matrices

(n,p)	$I(n, p)$	Total Matrices (p^{n^2})	$P(\text{irreducible})$	$P(\text{reducible})$
(2,2)	4	16	0.2500	0.7500
(2,3)	21	81	0.2593	0.7407
(3,2)	50	512	0.0977	0.9023
(3,3)	3459	19683	0.1757	0.8243

Table 3: Probability of a Reducible Minimal Polynomial

$p \setminus n$	2	3	4
2	0.7500	0.9023	0.9367
3	0.7407	0.8243	0.8729
5	0.6720	0.7542	0.8171

Asymptotic Comment. We analyze the behavior of the probability

$$P(n, p) = \frac{I(n, p)}{p^{n^2}}$$

that a random matrix in $M_n(\mathbb{F}_p)$ has an irreducible minimal polynomial. We have:

$$P(n, p) = \sum_{d|n} \frac{1}{p^{n^2}} \cdot N_d(p) \cdot \frac{|\text{GL}_n(\mathbb{F}_p)|}{|\text{GL}_{n/d}(\mathbb{F}_{p^d})|}$$

where

$$N_d(p) = \frac{1}{d} \sum_{k|d} \mu\left(\frac{d}{k}\right) p^k \quad \text{and} \quad |\text{GL}_m(\mathbb{F}_q)| = \prod_{i=0}^{m-1} (q^m - q^i).$$

Fix n . This is a finite sum over the divisors $d \mid n$, and we can analyze each term asymptotically. For such a divisor d , as $p \rightarrow \infty$, we have:

$$N_d(p) = \frac{p^d}{d} + \mathcal{O}(p^{d/2}),$$

and

$$\frac{|\text{GL}_n(\mathbb{F}_p)|}{|\text{GL}_{n/d}(\mathbb{F}_{p^d})|} = \frac{p^{n^2(1+o(1))}}{p^{n^2/d(1+o(1))}} = p^{n^2(1-\frac{1}{d})}(1+o(1)).$$

So the term corresponding to d in $P(n, p)$ behaves as

$$\frac{p^d}{d} \cdot p^{n^2(1-\frac{1}{d})} \cdot p^{-n^2}(1+o(1)) = \frac{1}{d} p^{d-\frac{n^2}{d}}(1+o(1)).$$

This tends to 0 as $p \rightarrow \infty$ if $d < n$, and to $\frac{1}{n}$ if $d = n$. Hence,

$$P(n, p) \longrightarrow \frac{1}{n} \quad \text{as } p \rightarrow \infty.$$

Moreover, based on initial values from a table, it appears that $P(n, p)$ increases with p . Therefore, we may conjecture that for all n and p ,

$$P(n, p) \leq \frac{1}{n} \leq \frac{1}{2}.$$

References

- [1] J. E. Cremona and D. Rusin. “Efficient solution of rational conics”. In: *Mathematics of Computation* 72.243 (2003), pp. 1417–1441.
- [2] Alexander J. Diesl, Samuel J. Dittmer, and Pace P. Nielsen. “Idempotent lifting and ring extensions”. In: *Proceedings of the American Mathematical Society* 143.9 (2015), pp. 3807–3811. DOI: [10.1090/S0002-9939-2015-12543-9](https://doi.org/10.1090/S0002-9939-2015-12543-9).
- [3] W. Eberly and M. Giesbrecht. “Efficient decomposition of associative algebras”. In: *Proceedings of the 1996 International Symposium on Symbolic and Algebraic Computation (ISSAC’96)*. Ed. by Y. N. Lakshman. New York: ACM, 1996, pp. 170–178.
- [4] Wayne Eberly. “Computations for Algebras and Group Representations”. <http://www.cpsc.ucalgary.ca/~eberly/Research/Papers/phdthesis.pdf>. PhD thesis. University of Toronto, 1989.
- [5] Philippe Gille and Tamás Szamuely. *Central Simple Algebras and Galois Cohomology*. Vol. 101. Cambridge Studies in Advanced Mathematics. Cambridge University Press, 2006. ISBN: 978-0-521-86103-8. DOI: [10.1017/CB09780511607219](https://doi.org/10.1017/CB09780511607219).
- [6] Gábor Ivanyos, Ádám D. Lekes, and Lajos Rónyai. *Improved algorithms for splitting full matrix algebras*. 2012. arXiv: [1211.1356](https://arxiv.org/abs/1211.1356) [math.RA]. URL: <https://arxiv.org/abs/1211.1356>.
- [7] Gábor Ivanyos and Lajos Rónyai. “Finding maximal orders in semisimple algebras over \mathbb{Q} ”. In: *Computational Complexity* 3 (1993), pp. 245–261. DOI: [10.1007/BF01271370](https://doi.org/10.1007/BF01271370). URL: <https://link.springer.com/article/10.1007/BF01271370>.
- [8] Gábor Ivanyos and Ágnes Szántó. “Lattice basis reduction for indefinite forms and an application”. In: *Discrete Mathematics* 153.1-3 (1996). Proceedings of the 5th Conference on Formal Power Series and Algebraic Combinatorics (Florence, 1993), pp. 177–188.
- [9] Lajos Rónyai. “Computing the structure of finite algebras”. In: *Journal of Symbolic Computation* 9.3 (1990), pp. 355–373. DOI: [10.1016/S0747-7171\(08\)80072-7](https://doi.org/10.1016/S0747-7171(08)80072-7).
- [10] Denis Simon. “Quadratic equations in dimensions 4, 5 and more.” In: *LMNO - UMR 6139, Université de Caen – France* (May 2005). 11th May 2005. URL: <mailto:simon@math.unicaen.fr>.

- [11] Denis Simon. “Solving norm equations in relative number fields using S -units”. In: *Mathematics of Computation* 71.239 (2002), pp. 1301–1321. ISSN: 0025-5718. DOI: [10.1090/S0025-5718-02-01309-1](https://doi.org/10.1090/S0025-5718-02-01309-1). URL: <https://www.ams.org/journals/mcom/2002-71-239/S0025-5718-02-01309-1/S0025-5718-02-01309-1.pdf>.
- [12] Denis Simon. “Solving quadratic equations using reduced unimodular quadratic forms”. In: *Mathematics of Computation* 74.251 (2005), pp. 1531–1543.
- [13] John Voight. *Identifying the matrix ring: algorithms for quaternion algebras and quadratic forms*. 2012. arXiv: [1004.0994](https://arxiv.org/abs/1004.0994) [math.NT]. URL: <https://arxiv.org/abs/1004.0994>.
- [14] John Voight. *Quaternion Algebras*. Vol. 288. Graduate Texts in Mathematics. Cham: Springer, 2021. ISBN: 978-3-030-56692-0. DOI: [10.1007/978-3-030-56694-4](https://doi.org/10.1007/978-3-030-56694-4). URL: <https://link.springer.com/book/10.1007/978-3-030-56694-4>.